

# Contents

<b>Message Passing Interface</b>	<b>1</b>
Services . . . . .	2
Program Organisation . . . . .	2
Starting and Finishing . . . . .	2
<b>Messages</b>	<b>2</b>
Basics . . . . .	3
Message Protocol . . . . .	3
Send and Receive Synchronisation . . . . .	5
Broadcasting a Message . . . . .	6
Reduce . . . . .	7
Scatter . . . . .	7
Gather . . . . .	7
Data Packaging . . . . .	8
<b>Timing Programs</b>	<b>8</b>
<b>Compiling and Running MPI</b>	<b>8</b>

## Message Passing Interface

- Standard message-passing library
  - Includes best of several previous libraries
- Versions for C/C++ and FORTRAN
- Available for free
- Can be installed on
  - Networks for workstations
  - Parallel computers
- There is no shared memory in MPI
- Each MPI process has its own separate address space
  - A pointer from one address space is useless in another
- Aimed at distributed memory machines

- Giant supercomputers made of clusters
  - Machines in separate locations
- But it can be implemented on a physically shared-memory machine

## Services

- Hide details of architecture
- Hide details of message passing, buffering
- Provides message management services
  - packaging
  - send, receive
  - broadcast, reduce, scatter, gather message modes

## Program Organisation

- SPMD - Single Program, Multiple Data
  - Every processor runs the same program
  - Each processor computes with different data
  - Variation of computation on different processors through if or switch statements
- MIMD in a SPMD framework
  - Different processors can follow different computation paths
  - Branch on if or switch based on processor identity

## Starting and Finishing

- Statement needed in every program before any other MPI code: `MPI_Init(&argc, &argv)`
- Last statement of MPI code must be: `MPI_Finalize()`
  - Program will not terminate without this statement

## Messages

- Message content, a sequence of bytes
- Message needs wrapper
  - Analogous to an envelope for a letter

Letter	Message
Address	Destination
Return Address	Source
Type of Mailing (class)	Message type
Letter Weight	Size (count)
Country	Communicator
Magazine	Broadcast

## Basics

### Communicator

- Collection of processes
- Determines scope to which messages are relative
- Identity of process (rank) is relative to communicator
- Scope of global communications (broadcast, etc.)

## Message Protocol

### Send

- Message contents: block of memory
- Count: number of items in a message
- Message type: type of each item
- Destination: rank of processor to receive
- Tag: integer designator for message
- Communicator: the communicator within which the message is sent

### Receive

- Message contents: buffer in memory to store received message
- Count: number of items in message
- Message type: type of each item
- Source: rank of processor sending
- Tag: integer designator for message
- Communicator: the communicator within which the message is sent
- Status: information about message received

Example

```
#include <stdio.h>
#include <string.h>
#include "mpi.h"      /* includes MPI library code specs */

#define MAXSIZE 100

int main(int argc, char* argv[])
{
    int myRank;          /* rank (identity) of process */
    int numProc;         /* number of processors */
    int source;          /* rank of sender */
    int dest;            /* rank of destination */
    int tag = 0;         /* tag to distinguish messages */
    char mess[MAXSIZE]; /* message (other types possible) */
    int count;           /* number of items in message */
    MPI_Status status;   /* status of message received */

    MPI_Init(&argc, &argv);          /* start MPI */

    /* get number of processes */
    MPI_Comm_size(MPI_COMM_WORLD, &numProc);

    /* get rank of this process */
    MPI_Comm_rank(MPI_COMM_WORLD, &myRank);

    if (myRank != 0) { /* all processes send to root */
        /* create message */
        sprintf(mess, "Hello from %d", myRank);
        dest = 0;          /* destination is root */
        count = strlen(mess) + 1; /* include '\0' in message */

        MPI_Send(mess, count, MPI_CHAR, 0, tag, MPI_COMM_WORLD);
    }
    else { /* root (0) process receives and prints messages */
        /* from each processor in rank order */
        for(source = 1; source < numProc; source++){
            MPI_Recv(mess, MAXSIZE, MPI_CHAR,
                     source, tag, MPI_COMM_WORLD, &status);
            printf("%s\n", mess);
        }
    }

    MPI_Finalize();          /* shut down MPI */
}
```

- Send - Receive is point-to-point, destination processor is specified by fourth parameter (**dest**) in **MPI\_Send**
- Messages can be tagged by integer to distinguish messages with different purposes by the fifth argument in **MPI\_Send** and **MPI\_Recv**
- **MPI\_Recv** can specify a specific source from which to receive (fourth parameter)
- **MPI\_Recv** can receive from any source or with any tag using **MPI\_ANY\_SOURCE** and **MPI\_ANY\_TAG**
- Communicator, switch parameter in **MPI\_Send** and **MPI\_Recv** determines context for destination and source ranks
- **MPI\_COMM\_WORLD** is automatically supplied communicator which includes all processes created at start-up
- Other communicators can be defined by use to group processes and to create virtual topologies
- Status of message received by **MPI\_Recv** is returned in the seventh (status) parameter
- Number of items actually received can be determined from status by using function **MPI\_Get\_count**
- The following call inserted into the previous code would return the number of characters sent in the integer variable **count**: **MPI\_Get\_count(&status, MPI\_CHAR, &count);**

### Send and Receive Synchronisation

- Fully Synchronised (Rendezvous)
  - Send and Receive complete simultaneously
    - \* Whichever code reaches the send/receive first waits
  - Provides synchronisation point (up to network delays)
- Buffered
  - Receive must wait until message is received
  - Send completes when message is moved to buffer clearing memory of message for reuse
- Asynchronous (different API call)
  - Sending process may proceed immediately
    - \* Does not need to wait until message is copied to buffer
    - \* Must check for completion before using message memory
  - Receiving process may proceed immediately
    - \* Will not have message to use until it is received
    - \* Must check for completion before using message

`MPI_Send` may be fully synchronous or may be buffered

`MPI_Recv` suspends until message is received

`MPI_Isend/MPI_Irecv()` are non-blocked. Control returns to program after call is made. (Syntax is the same except `MPI_Request` replaced `MPI_Status`).

Detecting completion

- `MPI_Wait(&request, &status)`
  - `request`: matches request on `Isend` or `Irecv`
  - `status`: returns status equivalent to status for `Recv` when complete
  - Blocks for send until message is buffered or sent so message variable is free
  - Blocks for receive until message is received and ready
- `MPI_Test(&request, flag, &status)`
  - `request, status` as for `MPI_Wait`
  - Does not block
  - Flag indicates whether message is sent/received
  - Enables code which can repeatedly check for communication completion

### Broadcasting a Message

- Broadcast: one sender, many receivers
- Includes all processes in communicator, all processes must make an equivalent call to `MPI_Bcast`
- Any processor may be sender (root), as determined by the fourth parameter
- First three parameters specify message as for `MPI_Send` and `MPI_Recv`, fifth parameter specifies communicator
- Broadcast serves as a global synchronisation

`MPI_Bcast(mess, count, MPI_INT, root, MPI_COMM_WORLD);`

- `mess`: pointer to message buffer
- `count`: number of items sent
- `MPI_INT`: type of item sent
- `root`: sending processor
- `MPI_COMM_WORLD`: communicator within which broadcast takes place

`MPI_Barrier(MPI_COMM_WORLD)`

- Provides for barrier synchronisation without message of broadcast
- A barrier is a point in the code where all processes must stop and wait until every process has reached the barrier
- Once all processes have executed the `MPI_Barrier` call then all processes can continue

## Reduce

- All processors send to a single processor, the reverse of broadcast
- Information must be combined at receiver
- Several combining functions available
- Data and result may be arrays - combining operation applied element-by-element
- Illegal to alias `dataIn` and `result`

```
MPI_Reduce(&dataIn, &result, count, MPI_DOUBLE, MPI_SUM, root,
MPI_COMM_WORLD)
```

- `dataIn`: data send from each processor
- `result`: store result of combining operation
- `count`: number of items in each `dataIn` result
- `MPI_DOUBLE`: data type for `dataIn`
- `MPI_SUM`: combining operation
- `root`: rank of processor receiving data
- `MPI_COMM_WORLD`: communicator

## Scatter

```
MPI_Scatter()
```

- Spreads array to all processors
- Source is an array on the sending processor
- Each receiver, including sender, gets a piece of the array corresponding to their rank in the communicator

## Gather

- `MPI_Gather()`
- Opposite of scatter
- Values on all processors (in the communicator) are collected into an array on the receiver
- Array locations correspond to ranks of processors

## Data Packaging

- Needed to combine irregular, non-contiguous data into single message
- pack/unpack, explicitly pack data into a buffer, send, unpack data from buffer
- Derived data types, MPI heterogeneous data types which can be sent as a message

`MPI_Pack(Aptr, count, MPI_DOUBLE, buffer, size, &pos, MPI_COMM_WORLD)`

- `Aptr`: pointer to data to pack
- `count`: number of items to pack
- `buffer`: buffer being packed
- `size`: size of buffer (in bytes)
- `pos`: position in buffer (in bytes), updated communicator

`MPI_Unpack(buffer, size, &pos, Aptr, count, MPI_DOUBLE, MPI_COMM_WORLD)`

## Timing Programs

- `MPI_Wtime()`
  - returns a double giving time in seconds from a fixed time in the past
  - To time a program, record `MPI_Wtime()` in a variable at start, then again at finish, difference is elapsed time

## Compiling and Running MPI

- MPI programs are compiled with a standard compiler such as `gcc`
- However, the compiler takes a lot of command line parameters to include the correct files and pass the right flags
- MPI implementations usually provide a special version of the compiler just for MPI programs
- In OpenMPI this is `mpicc`
  - `mpicc program.c -o program`
  - Internally this works by calling `gcc` with all the necessary flags and includes
- In OpenMPI you need to provide a command-line parameter to say how many copies of your program you want to run
  - `mpirun -np 10 ./program`



- Running MPI programs also requires a lot of flags and environment variables
  - Dynamic linking with MPI run-time libraries
- In OpenMPI a special program is used to load, dynamically link, and run your compiled MPI program