

Contents

Instruction Scheduling	1
Dynamic Instruction Scheduling	1
Dependence Types	2
Register Renaming	2
Simplified	2
Dynamic Instruction Scheduling Mechanisms	3
Content Addressable Memory	3
Out-of-order Benefits and Challenges	4

Instruction Scheduling

- Scheduling: act of finding independent instructions
 - “Static” done at compile time by the compiler (software)
 - “Dynamic” done at runtime by the processor (hardware)
- Why schedule code?
 - Scalar pipelines: fill in load-to-use delay slots to improve CPI
 - Superscalar: place independent instructions together
 - * As above, load-to-use delay slots
 - * Allow multiple-issue decode logic to let them execute at the same time

Dynamic Instruction Scheduling

- Also called “out-of-order” processors
- Hardware re-schedules instructions within a sliding window of instructions
- As with pipelining and superscalar, ISA unchanged
 - Same hardware/software interface, appearance of in-order
- Increases scheduling scope
 - Does loop unrolling transparently
 - Uses branch prediction to “unroll” branches
- “Dynamic scheduling” done by hardware

- Still 2-wide superscalar, but now out-of-order too
 - Allows instructions to issue when dependencies are ready
- Longer pipeline
 - In-order front end: Fetch, “Dispatch”
 - Out-of-order execution core: “Issue”, “RegisterRead”, Execute, Memory, Writeback
 - In-order retirement: “Commit”

Dependence Types

- **RAW** (read after write): true dependence
- **WAW** (write after write): output dependence
- **WAR** (write after read): anti-dependence

WAW & WAR are false; can be totally eliminated by renaming

Register Renaming

- To eliminate register conflicts/hazards
- Architectural vs Physical registers - level of indirection
- Renaming - conceptually write each register once
 - Removes false dependencies
 - Leaves true dependencies intact
- When to reuse a physical register?
 - After overwriting instructions is complete

Simplified

- Two key data structures:
 - `mactable[architectural_reg] -> physical_reg`
 - Free list: allocate (new) and free registers (implemented as a queue)
 - * Ignore freeing of registers for now
- Algorithm: at decode stage for each instruction
 - Rewrites instruction with physical registers (rather than architectural registers)

Dynamic Instruction Scheduling Mechanisms

Dispatch

- Put renamed instructions into out-of-order structures
- Re-order buffer
 - Holds instructions from Fetch through Commit
- Issue Queue
 - Central piece of scheduling logic
 - Holds instructions from Dispatch through Issue
 - Tracks ready inputs
 - * Physical register names + ready bit
 - * “AND” the bits to tell is ready

Out-of-order Pipeline

- Execution (out-of-order) stages
- **Select** ready instructions
 - Send for execution
- **Wakeup** dependents

Content Addressable Memory

- A content addressable memory (CAM) is indexed by the **content** of each location, not by the address
 - Sometimes known as associative memory
- It compares an input key against a table of keys, and returns the location of the key in the table
 - In software this might be implemented with a hash table
 - Hardware hash table is also possible, but potentially slow
- To search all locations in a single cycle
 - You need to be able to compare the search key to all keys in the table simultaneously
 - * This requires a lot of hardware
 - * Fast CAMs are very hardware expensive
 - If you need to be able to do multiple searches in the same cycle, the hardware requirements are even greater

Out-of-order Benefits and Challenges

Benefits

- Allows speculative re-ordering
 - Loads/stores
 - Branch prediction to look past branches
- Done by hardware
 - Compiler may want different schedule for different hardware configs
 - Hardware has only its own configuration to deal with
- Schedule can change due to cache misses
- Memory-level parallelism
 - Executes “around” cache misses to find independent instructions
 - Finds and initiates independent misses, reducing memory latency

Challenges

- Design complexity
 - More complicated than in-order? Certainly!
 - But we have managed to overcome the design complexity
- Block frequency
 - Can we build a “high ILP” machine at a high clock frequency?
 - Yep with some additional pipe stages, clever design
- Limits to (efficiently) scaling the window and ILP
 - Large physical register file
 - Fast register renaming/wakeup/select/load queue/store queue
 - * Active areas of micro-architectural research
 - Branch & memory depend prediction (limits effective window size)
 - * 95% branch misprediction: 1 in 20 branches, or 1 in 100 insns
 - Plus all the issues of building wide in-order superscalar
- Power efficiency
 - Today, even mobile phone chips are out-of-order cores