

Contents

Relational Algebra	1
Operations	2
Selection	2
Selection in SQL	5
Projection	5
Projection in SQL	6
Duplicate Elimination	6
Combining Select and Project	6
Nesting Operations	6
Immediate Relations	7
Combining Select and Project	7
Set Operators	7
Union	8
Intersection	8
Union and Intersection	8
Difference	8
Set Operators in SQL	8
Join	9
Joins in SQL	9

Relational Algebra

- A data model must include a set of operations to manipulate the database
 - Relational Algebra is the basic set of operations for the Relational Model
- Relational Algebra is important, as it:
 - Provides a formal foundation for relational model operations
 - Is used as a basis for implementing and optimising queries
 - Is incorporated into the SQL language for RDBMS

- Relational Algebra is a collection of operations on relations which fall into two groups
 - Set operations from mathematical set theory
 - * Union, intersection, etc.
 - * Applicable as each relation is defined to be a set of tuples in the relational model
 - Relational operations
 - * Defined specifically for relational databases
 - * Selection, projection, join, etc.

Operations

- Set operations
 - Standard mathematical operations on sets
 - UNION, INTERSECTION, SET DIFFERENCE
- Relational Database Operations
 - Unary Operations
 - * Operate on a single relation
 - * SELECT and PROJECT
 - Binary Operations
 - * A JOIN is used to combine related tuples across two relations using a *join condition*

Selection

- The SELECT operation is used to identify the subset of tuples from a relation that satisfy a selection condition
 - Acts as a filter on a relation
 - Horizontal partition of a relation
- Formal Notation
- $\sigma_{(selection\ condition)}(R)$
- So to identify all the employees in department 4

EMPLOYEE

Fname	Minit	Lname	Ssn	Dno
John	B	Smith	123456789	5

Fname	Minit	Lname	Ssn	Dno
Franklin	T	Wong	333445555	5
Alicia	J	Zelaya	999887777	4
Jennier	S	Wallace	Z987654321	4

- $\sigma_{(Dno=4)}$ (EMPLOYEE)
- The result of SELECT operation is a new relation
 - Made up of those tuples that satisfied the selection condition

Fname	Minit	Lname	Ssn	Dno
Alicia	J	Zelaya	999887777	4
Jennier	S	Wallace	Z987654321	4

- This new relation has the same attributes as the relation **R**, upon which the selection was performed
- Boolean Expression specified in the (selection condition) is made up of a number of clauses:
 - `<attribute name><comparison op><constant value>`
- `<attribute name>`
 - name of an attribute of **R**
- `<comparison op>`
 - operator such as `= < ≤ > ≥ ≠`
- `<constant value>`
 - a constant value from the attribute domain
- Selection conditions can be joined by boolean operators *and*, *or* and *not*
 - (cond1 AND cond2) is TRUE if both (cond1) and (cond2) are TRUE; otherwise, it is FALSE
 - (cond1 OR cond2) is TRUE if either (cond1) or (cond2) or both are TRUE; otherwise, it is FALSE
 - (NOT cond) is TRUE if cond is FALSE; otherwise, it is FALSE
 - Example
 - * Select all employees who either work in department 4 and make over €25,000 a year or work in department 5 and make over €30,000 a year

- * $\delta_{(Dno=4 \wedge Salary > 25000)} \text{ or } \delta_{(Dno=5 \wedge Salary > 30000)}(\text{EMPLOYEE})$
- Rather than using a constant value to evaluate the selection condition, another attribute can be used
 -
 - Example
 - * Select all employees who earn more than the lower tax band
 - * $\delta_{(Salary > lower_tax_band)}(\text{EMPLOYEE})$
- Full set of comparison operators ($= < \leq > \geq \neq$) can be used on all *Ordered Domains*
 - Numeric, Currency, Dates
- Can also have *Unordered Domains*
 - i.e. Color = {Red, Blue, Green, Yellow, ...}
 - Only valid comparison operators
 - * $\{=, \neq\}$
 - * Exceptions - substring
- The Select operator is Unary
 - Applied to a single relation
 - Select is applied to each tuple in turn
- Degree
 - The degree is the number of attributes in the relation
 - The degree of the Select operator is the same as the degree of the relation R
- The number of tuples returned for a Select operation, is always less than or equal to the number of tuples in R
 - $|\delta_c(R)| \leq |R|$
- The fraction of tuples in a relation selected by a condition is known as the *selectivity* of that condition
- The Select operation is *commutative*
 - $\delta_{(condition1)}(\delta_{condition2}(R))$
 - $\delta_{(condition2)}(\delta_{condition1}(R))$
- A “*cascade*” or sequence of select operations can be combined into a single operation
 - $\delta_{(condition1)}(\delta_{condition2}(\dots(\delta_{(conditionN)})\dots))$
 - $\delta_{(condition1)} \wedge (condition2) \wedge \dots \wedge (conditionN)(R)$

Selection in SQL

- SELECT is a common command in SQL, directly based upon this Relational Algebra approach
- While the algebraic notation is:
 - $\delta_{Dno=4 \wedge Salary > 25000}(EMPLOYEE)$
- The SQL syntax is:

```
SELECT * from EMPLOYEE
WHERE Dno=4
AND Salary > 25,000;
```

Projection

- The Project operation selects certain attributes from the table, while discarding the others
 - Vertical partition of relation
- Formal Notation
 - $\pi_{attribute\ list}(R)$
- Example
 - To list all employees first and last names and their salary, we can use Project:
 - $\pi_{Fname,Lname,Salary}(EMPLOYEE)$
- The relation that is the result of the Project operation has only the attributes specified in the attribute list
 - Hence, the *degree* of the resulting relation is equal to the number of attributes in this list
- Duplicate Elimination
 - It is possible for a Project operation to specify a set of non-key attributes, i.e. forename and surname
 - In this instance, duplicate tuples are likely to occur
 - However, the result of a Project operation has to be a valid relation, and only contain distinct tuples
 - The operation removes (or merges) the duplicate tuples and only returns distinct tuples

Projection in SQL

- The SELECT statement in SQL also handles the Project operation
- While the algebraic notation is
 - $\pi_{Fname, Lname, Salary}(EMPLOYEE)$
- The SQL syntax is:

```
SELECT Fname, Lname, Salary
from EMPLOYEE;
```

Duplicate Elimination

- SQL does not automatically implement the duplicate elimination process of the Project operation
 - This is a departure from formal Relational Algebra
- To implement this, you need to use the DISTINCT keyword

```
SELECT DISTINCT firstname, surname, salary
from EMPLOYEE;
```

Combining Select and Project

- Complex queries are often need to use combinations of Select and Project operations
- This can be achieved in two ways
 - Nesting of operations in a single relational algebra expression
 - One operation at a time using intermediate result relations

Nesting Operations

- If we want a list of the names and salaries of all employees who earn more than €25,000 and work in the administration department
- This would require two operations
 - $\delta_{Dno=4 \wedge Salary > 25000}(EMPLOYEE)$
 - $\pi_{Fname, Lname, Salary}(EMPLOYEE)$
- However, we can achieve the same result in a single relational algebra expression by nesting
 - $\pi_{Fname, Lname, Salary}(\delta_{Dno=4 \wedge Salary > 25000}(EMPLOYEE))$

Immediate Relations

- Operations can be applied one at a time using intermediate relations to store the results
- All intermediate tables must be named so that they can be referred to
 - $EMPS \leftarrow \delta_{(Dno=4 \wedge Salary > 25000)}(EMPLOYEE)$
 - $RESULTS \leftarrow \pi_{Fname, Lname, Salary}(EMPS)$
- It can sometimes be simpler to break complex queries into a series in this manner

Combining Select and Project

- The Select and Project operations can be combined in the SQL SELECT statement
- Consider the following operations
 - $\delta_{Dno=4 \wedge Salary > 25000}(EMPLOYEE)$
 - $\pi_{Fname, Lname, Salary}(EMPLOYEE)$
- These can be combined in the following SQL

```
SELECT Fname, Lname, Salary
from EMPLOYEE
WHERE Dno=4
AND Salary>25,000
```

Set Operators

- Standard mathematical operations used to merge the elements of two Sets
 - Union
 - Intersection
 - Set Difference
- Binary Operations
 - Each operation is applied to two sets
- To use these operations, two relations must be “union compatible”
 - They must be of the same number of attributes
 - Each corresponding pair of attributes has the same domain
- Two relations $R(A_1, A_2, \dots, A_n)$ and $S(B_1, B_2, \dots, B_n)$ are said to be union compatible if
 - They have the same degree n
 - $dom(A_i) = dom(B_i)$ for $1 \leq i \leq n$

Union

- Denoted $R \cup S$
- The result of a *Union* operation is a new relation which contains all tuples that are either in **R** or **S** or in both **R** and **S**
- Duplicate tuples are discarded
- Commutative operation
 - $R \cup S = S \cup R$

Intersection

- Denotes $R \cap S$
- The result of an *Intersection* operation is a new relation which contains all tuples that are in both **R** and **S**
- Commutative operation
 - $R \cap S = S \cap R$

Union and Intersection

- Can both be treated as **n-ary** operations
 - Applicable to any number of relations
 - As they are *associative operations*
 - $R \cup (S \cup T) = (R \cup S) \cup T$
 - $R \cap (S \cap T) = (R \cap S) \cap T$

Difference

- Denoted by **R - S**
- The result of a *Set Difference* (or *Minus*) operation is a new relation that includes all tuples that are in **R** but not in **S**
- Not commutative
 - $R - S \neq S - R$

Set Operators in SQL

- There are three SQL commands which correspond to these Set Operations
 - **UNION** - Union
 - **INTERSECT** - Intersection
 - **EXCEPT** - Set Difference

- Additionally, there are multi-set operators in SQL that do not eliminate duplicates
 - UNION ALL
 - INTERSECT ALL
 - EXCEPT ALL

Join

- Denoted by $R \bowtie_{join\ condition} S$
- Used to combine related tuples from two relations, into a single tuple
 - Important as it allows for the processing of relationships between relations
- Makes use of Foreign Keys and Referential Integrity constraints that have been defined
- Consider two relations $R(A_1, A_2, \dots, A_n)$ and $S(B_1, B_2, \dots, B_m)$
- The result of $R \bowtie S$ is a new relation Q with
 - $n+m$ attributes
 - $Q(A_1, A_2, \dots, A_n, B_1, B_2, \dots, B_m)$
 - A tuple for each combination of tuples (one from R and one from S) that satisfy the join condition
- The join condition is specified on attributes from both relations, R and S
 - It is evaluated for every combination of tuples from the two relations
 - Each tuple combination for which the join condition is TRUE is included in the result relation as a single, combined tuple
- Suppose we want to return the name of the manager of each department
 - $DEPARTMENT \bowtie_{Mgr_sn=Ssn} EMPLOYEE$
- Can then be combined with a Project operation to only return the required attributes
 - $\pi_{Dname, Fname, Lname}(DEPARTMENT \bowtie_{Mgr_sn=Ssn} EMPLOYEE)$

Joins in SQL

```
SELECT employee.name, job, department.name
FROM employee, department
WHERE employee.deptno = department.deptno;
```

or

```
SELECT employee.name, job, department.name  
FROM employee  
INNER JOIN department  
ON employee.deptno = department.deptno
```