

Contents

Key Points about Relational Models	2
Duplicated vs Redundant Data	3
Example	3
Eliminating Redundancy	3
Repeated Groups	4
Problems	5
Elimination of Repeating Groups	5
Alternative Method	5
Eliminating Repeating groups and Redundancy	6
Conclusion	6
Codd's Normal Forms	6
Determinants	7
Example	7
Superfluous Attribute	7
Diagrams	8
Transitive Determinants	8
Identifiers	8
Example	8
Rules	8
Determinancy and Redundancy	8
Transforming Tables in Well-Normalised Table	9
Fully Normalised Tables	9
Example	10
Advantages of Full Normalisation	10

Conceptual Design	11
Entity Relationship Model	11
Entities	12
Attributes	12
Simple and Composition Attributes	12
Single and Multi-Valued Attributes	12
Stored versus Derived Attributes	13
Entity Types and Sets	13
Key Attributes	13
Relationships	14
Relationship Types and Sets	14
Binary and Ternary Relationships	15
Relationship Roles	15
Recursive Relationships	15
Relationship Constraints	15
Cardinality Constraints	15
Participation Constraints	16
Relationship Type Attributes	16
Relationships as Attributes	17
Weak Entity Types	17

Key Points about Relational Models

1. Ordering of rows is not significant
2. Ordering of columns is not significant
3. Each row/column intersection contains a single attribute value
4. Each row in a table must be distinct

A table conforming to these restrictions is called a *normalised* table

Duplicated vs Redundant Data

- Must be careful to distinguish between redundant and duplicated data
- Duplicated data occurs where an attribute (column) has two or more identical values
- Redundant data occurs if you can delete a value without information being lost

Redundancy is unnecessary duplication

Example

Part#	Part Desc
P1	bolt
P2	nut
P3	washer
P4	nut

nut is duplicated but **NOT** redundant

Part#	Part Desc
P1	bolt
P2	nut
P3	washer
P1	nut

nut is duplicated and redundant. No loss in information.

Eliminating Redundancy

- We cannot just delete values from the table in the previous example!
- Preferable to split table into 2 tables - parts and suppliers

Part#	Part Desc
P1	bolt
P6	bolt
P4	nut

S#	Part#
S2	P1
S7	P6
S2	P4
S5	P1

- Eliminated redundancy by table splitting
 - P1 description only appears once
 - Relationship is made by including part # in two tables
- So far we've assumed that table structures which permit redundancy can be recognised by inspection of table occurrence
- This is *not entirely accurate* since attribute values are subject to insertion/change/deletion

S#	P#	Desc
S2	P1	bolt
S7	P6	bolt
S2	P4	nut

Inspection of table does not reveal any redundancy

Could even suggest that “no two suppliers may supply same part #”

Repeated Groups

- We stated earlier that “an attribute must only have one value in each row”

S#	SName	P#
S5	Wells	P1
S2	Heath	P1, P4
S7	Barron	P6
S9	Edwards	P8, P2, P6

Problems

1. Table is asymmetric representation of symmetrical data
2. Rows can be sorted into s# but not into p#
3. Rows are different length because of variation in number of p#s
4. If rows were fixed length, they would need to be padded with null values

Elimination of Repeating Groups

- Easiest way to eliminate repeating groups is to write out the table occurrence using a vertical layout and fill in the blanks by duplicating the non-repeating data necessary
- But this can lead to 'redundancy' of information

S#	Sname	P#
S5	Wells	P1
S2	Heath	P1
S2	Heath	P4
S7	Barron	P6
S9	Edwards	P8
S9	Edwards	P2
S9	Edwards	P6

Alternative Method

- Split table into two tables so that repeating group appears in one table and rest of attributes in another
- Need to provide correspondance between tables by including a key attribute the with repeating group table

S#	SName
S5	Wells
S2	Heath
S7	Barron
S9	Edwards

S#	P#
S5	P1
S2	P1
S2	P4
S7	P6
S9	P8
S9	P2
S9	P6

Eliminating Repeating groups and Redundancy

- Snapshot of table is inadequate guide to presence/absence of redundant data
- Need to know underlying rules
- DBA must discover rules which apply to conceptual model

Conclusion

- Its not possible to tell by looking at the relational tables in a DB to determine if
 - There is the potential for redundancy
- But what would be a ‘correctly formed’ table?

Codd’s Normal Forms

- Codd identified some rules which govern the way we create tables so as to avoid anomalies when inserting or deleting values in these tables

- These rules are called *NORMAL forms*
 - There are three and a half important levels (and two further levels which are occasionally used)
1. A relational is in first normal form if the domain of each attribute contains only atomic values and the value of each attribute contains only a single value from that domain
 2. A relation is in second normal form if, in addition to satisfying the criteria for first normal form, every non-key column is *fully functionally dependent* on the entire primary key
 3. A relation is in the third normal form if, in the addition to satisfying the criteria for second normal form, and no non-key attributes are *transitively dependency* upon the primary key
- Boyce Codd Normal Form (Also called three and a half normal form)
 - “All attributes in a relation should be dependency on the key, the whole key and nothing but the key”

Determinants

- If there are rules such that duplicate values of attribute A are always associated with the *same* value of attribute B (within any given occurrence of the table) then attribute A is a *determinant* of attribute B

Note: In the special case where duplicate values of A are not allowed in a table (i.e. A is a key) then A is obviously a determinant on B

Example

- If each possible p# value has precisely one associated part description value (i.e. P4 has just one description nut then we can say that p# is a determinant of part description)
- Similarly, if each possible p# value has precisely one quantity in stock then we can say p# is a determinant of quantity in stock

Superfluous Attribute

- If P# determines Qty then composite attribute {P#, P_Desc} also determines Qty, but P_Desc is superfluous
- We assume determinants do not contain any superfluous attributes

Diagrams

- $A \rightarrow B$
 - A is determinant of B
- $A \leftrightarrow B$
 - A is determinant of B and vice versa

Transitive Determinants

- If A is determinant of B and B is determinant of C
- Then A is determinant of C

Identifiers

- Because of the rule that ‘no two rows in a table can have identical values throughout’
- Therefore individual row can always be identified by quoting the values of all its attributes
- However some values may not be needed

Example

`Employee(Employee#, Employee_name, Salary)`

Rules

- No two rows should have the same value for Employee#
 - Employee# is a row identifier of the table
- Where a composed attribute forms the identifier
 - No component part (of identifier) can be null (entity constraint)

Determinancy and Redundancy

- Given a determinancy diagram (developed from enterprise rules) we can detect and eliminate table structures which could contain redundant data

Customer# \rightarrow Salesman# \rightarrow Salesman name

Customer# is a determinant and identifier. Salesman# is a determinant only.

- Each customer# is associated with one salesman# but a salesman# may be associated with several different customer#
- Therefore salesman# could have duplicate values
- But salesman# is a determinant of salesman name
- Therefore each occurrence of a duplicate salesman# values will be associated with the same salesman name
 - Table can contain redundant values of salesman
- But customer# values cannot be duplicated (because customer# is the identifier for our table) so we cannot allow redundant values of salesman#
- Potential redundancy arises because salesman# is a determinant but not a candidate identifier

Transforming Tables in Well-Normalised Table

- Boyce Codd rule for determining redundancy is rule “Every determinant must be a candidate identifier”
- A table which obeys this rule is said to be in Boyce Codd normal form (BCNF)

To put it another way, “all attributes in a relation should be dependency on the key, the whole key and nothing but the key”

- A determinant which is not a candidate identifier is called a *non identifying determinant*
- To transform a badly normalised table into well normalised tables:
 - Create new tables such that each non identifying determinant in the old table becomes a candidate identifier in a new table

Fully Normalised Tables

- Fully normalised tables are structured in such a way that they cannot contain redundant data
- Generally a well normalised table (i.e. one in which each determinant is a candidate identifier) is also fully normalised, but not always!
 - So further normalisation may be desirable

Example

A database storing books has the following rules

- Each book has a unique book#
- Each author has a unique author#
- Every author has a name and every book has a title
- Each subject classification has a unique subject name
- Book# does not distinguish an individual copy of a book, only an individual work
- A book may be written by several authors, and be classified under several subject names
- An author may write several books
- A subject name may be applied to several books

AUTHOR-BOOK

Author#	Book#
A2	B15
A5	B15
A2	B18

BOOK-SUBJECT

Book#	Subject-Name
B15	Biology
B15	Physics
B18	Physics

Advantages of Full Normalisation

- So far emphasis has been placed on eliminating redundancy
- Further benefits relate to deletion, insertion operations

C#	CName	S#	SName
C1	Brown	S4	Jones

C#	CName	S#	SName
C2	Carter	S7	Samson
C3	Cross	S4	Jones
C4	Barns	S8	Baker

- Delete C2, delete whole tuple
 - Lose salesman information
- Deleting C# on its own is not allocated as it is an identifier and cannot be null

Insertion side effect

- Add S2 whose name is Hall
- You cannot do this until the salesman is associated with a customer, otherwise identifier C# will be null

Conceptual Design

- The process of *conceptual design* uses a high-level *conceptual data model* to create the *conceptual schema* for a database
 - In this case, the Relational Model
 - Part of the requirements analysis process
- Conceptual Schema
 - Concise description of the data requirements of the user, including descriptions of entity types, relationships and constraints
- The conceptual schema does not include any implementation details
 - Hence, can be used to communicate with non-technical users
 - Used as a reference to ensure all data requirements are met, and that there are no conflicts
- Part of physical and logical data independence

Entity Relationship Model

- An abstract and high-level conceptual representation of information

- Entity Relationship Diagrams
 - Diagrammatic Notation of the ER Model
- Used to support the conceptual design of databases and help produce the conceptual schema
- Describes data as entities, relationships and attributes

Entities

- The basic object that an ER diagram represents is an *entity*
- An entity is a real world object with an independent existence
 - Physical or conceptual
- Each entity has attributes which are the particular properties that describe the real world object

Attributes

- Several types of attributes occur and need to be modeled in the Entity-Relationship Model
 - *Simple* versus *composite*
 - *Single-valued* versus *multi-valued*
 - *Stored* versus *derived*

Simple and Composite Attributes

- Composite attributes can be divided into smaller sub-parts
 - Address
- Attributes that are not divisible are called simple, or atomic, attributes
 - Movie Certificate, Age...
- Composite attributes can be hierarchical
 - Apartment number, Building number, Street

Single and Multi-Valued Attributes

- Most attributes have a single attribute for each entity
 - PPS number, Age
 - Such attributes are called *single-valued*

- In some cases an attribute can have a set of values for an entity
 - Genre for a Movie, Colour for a Car
 - Some entities may have one value, others multiple
 - Such attributes are called multi-valued

Stored versus Derived Attributes

- In some cases, two or more attributes are related
 - Age and Birth Date
- Age can be calculated using today's date and a person's date of birth
 - Age is therefore called a derived attribute
 - It is said to be derivable from Birth Date
 - Birth Date is called a stored attribute
- Some attributes can be derived from information in related entities, rather than attributes
 - If a number of employees attribute was associated with a CINEMA entity
 - This could be derived by totaling the number of employee entries stored in the EMPLOYEE entity

Entity Types and Sets

- ER diagrams don't show single instances of entities (or relations)
- They show *entity types*
 - An entity which is identified by its name and attributes
 - In a cinema database, MOVIE could be an entity type
 - all movie entities share the same attributes but each instance has its own values for each attribute
- The collection of all instances of a particular entity type in a database is called an *entity set*

Key Attributes

- Each entity type usually has one or more attributes whose values are unique for each instance in the entity set
 - An attribute whose values are used to uniquely identify each entity is called the *key attribute*
 - * ISBN, PPS, Student Number

- More than one attribute can be used to form the key
 - In this case the combination must be unique
 - * Composite key attribute
- Specifying the key attribute places a uniqueness constraint on the entity type
 - This uniqueness constraint must hold for every instance of an entity in the entity set
- The key constraint is derived from the real world requirements that the database represents
- Some entity types have more than one key attribute

Relationships

- A *relationship* captures how two or more entity types are related to one another
- Whenever an attribute of an entity type refers to another entity type, a relationship exists
 - There are a number of implicit relationships in a movie example - SCREEN and THEATRE, SCREENING and MOVIE, SCREENING and SCREEN
 - In the ER model, these references should not be represented as attributes, but as relationships
- A Relationship can be informally thought of as a verb, linking two or more nouns from the world you are trying to model
 - A “manages” relationship between an employee and a department
 - A “performs” relationship between an artist and a song
 - A “bores” relationship between a lecturer and a student
 - A “proved” relationship between a mathematician and a theorem

Relationship Types and Sets

- As with entities, relationships have a *relationship type*, which is illustrated in an ER diagram
 - The collection of all instances of a particular relationship type in a database is called a *relationship set*
- Related entity types are said to *participate* in a relationship type
 - Each relationship instance, r_i , is an association of entities, where the association includes exactly one entity from each of the participating entity types

Binary and Ternary Relationships

- The degree of a relationship type is the number of entity types that participate
 - The SHOW relationship is of degree two
- Relationship types of degree two are called *binary*, relationship types of degree three are called *ternary*

Relationship Roles

- Each entity type that participates in a relationship type plays a particular role
- A role name can be optionally added to an ER diagram to clearly identify what the relationship means

Recursive Relationships

- In some cases the same entity type participates more than once in a relationship type, in different roles
 - Such relationships are called *recursive relationships*
- In this case the relationship roles are key in distinguishing the role each participating entity plays

Relationship Constraints

- Constraints limit the possible combination of entities that can participate in a relationship
- These constraints are determined by the real world requirements that are being modeled
- Two main types of relationship constraint
 - Cardinality constraints
 - Participation constraints

Cardinality Constraints

- Specify the *maximum number of relationship instances that an entity can participate in
- Consider our SHOW relationship type
 - Cardinality ratio for MOVIE:SCREENINGS is 1:N

- Each movie can be shown in many screenings
 - Each screening shows a maximum of one movie
- Possible cardinality ratios for binary relationships
 - 1:1 - One to one
 - 1:N - One to many
 - M:N - Many to many

Participation Constraints

- Specify the *minimum* number of relationship instances that an entity can participate in
- Two types of participation constraint
 - *Total participation*
 - *Partial participation*
- Formally, participation constraints specify if the existence of an entity depends on it being related to another entity via the relationship type
- Company policy states that every Employee MUST work for a department
- Then, an employee entity can only exist if they participate in at least one WORKS relationship between EMPLOYEE and DEPARTMENT
- The participation of EMPLOYEE in WORKS is called *total participation*
 - Every entity in the total set of employee entities must be related to a department entity via the WORKS relationship type
- Total participation is also called *existence dependency*
- In the same company, every employee will not be responsible for managing a department
- Therefore, an employee entity can exist if they don't participate in the MANAGE relationship between EMPLOYEE and DEPARTMENT
- The participation of EMPLOYEE in MANAGE is called *partial participation*
 - Some, but not all, of the total set of employee entities are related to a department entity via the MANAGE relationship type
- Together, cardinality constraints and participation constraints are referred to as the *structural constraints* of a relationship type
- In ER diagrams, total participation is represented by a double line connecting the participating entity type to the relationship type
- Partial participation is represented by a single line

Relationship Type Attributes

- In the entity relationship model, relationship types can have attributes, similar to entity types

- These relationship type attributes can be migrated to participating entities for relationships with cardinality of 1:1 or 1:N
- In relationship types with 1:N cardinality ratios, the relationship attribute can only be migrated to the N side of the relationship
- In relationship types with N:M cardinality ratios, the attribute cannot be migrated and remains specified as a relationship attribute

Relationships as Attributes

- It is sometimes convenient to think of a relationship type in terms of attributes
 - As when initially modelling the entity types
- With binary relationships there are always two options for representing it as an attribute
 - Consider the “shows” relationship
 - We could use an attribute in Screening that refers to the Movie ID for the movie in question
 - Or, we could use a multi-valued attribute in movie that refers to all the screenings of that movie

Weak Entity Types

- In ER diagrams, entity types which don’t have a key attribute are called *weak entity types*
- Entities of this type are identified by their relationship to specific entities from other entity types
 - This other entity type is called the identifying entity type
 - The relationship type that related a weak entity type to its identifier is called the identifying relationship
- A weak entity type always has a total participation constraints with relation to its identifying relationship
 - As a weak entity cannot be identified without its identifier entity