

# Contents

<b>Introduction</b>	<b>2</b>
<b>Transactions</b>	<b>2</b>
Transaction Failure . . . . .	3
Transaction Properties . . . . .	4
<b>Concurrency Control</b>	<b>4</b>
Lost Update . . . . .	5
Temporary Update . . . . .	5
Incorrect Summary . . . . .	5
<b>Schedules</b>	<b>5</b>
Schedule Conflicts . . . . .	6
Serial Schedules . . . . .	6
Issues with Serial Schedules . . . . .	7
Serializability . . . . .	7
Result Equivalence . . . . .	7
Conflict Equivalence . . . . .	7
Serializability . . . . .	8
<b>Concurrency Control Protocols</b>	<b>8</b>
<b>Locking</b>	<b>9</b>
Binary Lock . . . . .	9
Read/Write Lock . . . . .	9
Lock Conversion . . . . .	10
Two-Phase Locking . . . . .	10
Limitations . . . . .	11
Problems with Locking . . . . .	11
Deadlock . . . . .	12
Deadlock Prevention . . . . .	12
No Waiting . . . . .	12

Wait-die . . . . .	13
Wound-wait . . . . .	13
Wait-die and Would-wait . . . . .	13
Cautious Waiting . . . . .	13
Deadlock Detection . . . . .	14
Starvation . . . . .	14
Starvation Solution . . . . .	14
Timestamp Ordering . . . . .	15

## Introduction

- Airline reservations, Credit Card Processing, Online Shopping, etc. Are example of very large real-world database systems
  - Can potentially have thousands of users performing operations at the same time
  - Operations can be complex, involved a number of separate actions
    - \* Create booking
    - \* Update remaining seats
    - \* Charge credit card
    - \* Issue confirmation
    - \* ...
- *Transactions* and *Concurrency Control* are the ways in which a DB manages complex processes and multi-user access
- Transaction
  - A logical unit of DB processing that must be completed in its entirety to ensure correctness
- Concurrency Control
  - Used when two operations try to access the same data at the same time

## Transactions

- A transaction includes one or more DB access operations
  - Insertion
  - Deletion

- Modification
  - Retrieval
- Transactions where the DB operations retrieve data, but don't update any information
  - Read-Only Transactions
- Transactions which update the DB
  - Read-Write Transactions

```

BEGIN_TRANSACTION
    READ or WRITE operation1
    READ or WRITE operation2
    ...
    READ or WRITE operationN
END_TRANSACTION
COMMIT or ROLLBACK

```

- The end of a transaction is signalled by either
  - commit (successful termination)
  - rollback or abort (unsuccessful termination)
- Commit completes the current transaction, making its changes permanent
- Rollback will undo the operations in the current transaction, cancelling the changes made to the DB

## Transaction Failure

- There are a number of reasons why a transaction might fail
  - System crash
    - \* Hardware, Software or Network Error
  - Transaction Error
    - \* Divide by zero
    - \* Incorrect attribute reference
    - \* Incorrect parameter value
  - Check can be placed in the transaction
    - \* Insufficient funds to make a withdrawal

## Transaction Properties

- All transactions should possess the ACID properties
  - Atomicity
    - \* A transaction is an atomic unit of processing
    - \* It should either be performed in its entirety or not performed at all
  - Consistency Preservation
    - \* A transaction should preserve the consistency of the DB
    - \* It should take the database from one consistent state to another
  - Isolation
    - \* A transaction should appear as though it is being executed in isolation
    - \* The execution of a transaction should not be interfered with by any other transactions executed concurrently
  - Durability (or permanency)
    - \* The changes applied by a committed transaction must persist in the database
    - \* These changes must not be lost because of any failure

## Concurrency Control

- Consider the following two transactions

Flight Transfer

```
read(reserved_seats_X)
reserved_seats_X = reserved_seats_X-N
write(reserved_seats_X)
read(reserved_seats_Y)
reserved_seats_Y = reserved_seats_Y+N
write(reserved_seats_Y)
```

Flight Reservation

```
read(read(reserved_seats_X)
reserved_seats_X = reserved_seats_X+M
write(reserved_seats_X)
```

- A number of problems can occur if these two simple transactions are run concurrently

## Lost Update

- This occurs when two or more transactions:
  - Access the same data item
  - Are executed concurrently
  - Are interleaved in such a way that results in an incorrect value being written to the database

## Temporary Update

- This occurs when two or more transactions:
  - Access the same data item
  - Are executed concurrently
  - Are interleaved
  - One transaction fails and must rollback
- This is also known as the “Dirty Read”

## Incorrect Summary

- This occurs when two or more transactions:
  - Access the same data item
  - Are executed concurrently
  - Are interleaved
  - One transaction is calculating an aggregate summary on attributes while another transaction is updating those attributes

## Schedules

- When transactions are executing concurrently in an interleaved fashion, the order of execution of operations is called the *schedule*
  - Schedule S of n transactions  $T_1, T_2, \dots, T_n$  is the ordering of the operations within those transactions
- Operations from different transactions can be interleaved
  - The operations from each transaction  $T_i$  must appear in the same order in S, that they do in  $T_i$
- A schedule S is generally written:
  - $S_a: O_1(X); O_1(Y); O_2(X); \dots; O_n(Y);$

- Where  $O_i(X)$  indicates either a read or write operation executed by a transaction  $T_i$  on a data item  $X$
- And  $O_n(Y)$  indicates either a read or write operation executed by a transaction  $T_n$  on a data item  $Y$
- A shorthand notation can be used for each operation type
  - $b \rightarrow \text{BEGIN\_TRANSACTION}$
  - $r \rightarrow \text{Read Item}$
  - $w \rightarrow \text{Write Item}$
  - $e \rightarrow \text{END\_TRANSACTION}$
  - $c \rightarrow \text{Commit}$
  - $a \rightarrow \text{Abort (Rollback)}$

## Schedule Conflicts

- Two operations in a schedule are said to *conflict* if:
  1. They belong to *different transactions*
  2. They access the *same item X*
  3. And *at least one* of the operations is a write( $X$ )
- Intuitively, two operations are conflicting if changing their order can result in a different outcome
  - Or cause one of the concurrency issues we've already discussed

## Serial Schedules

- In a *serial schedule* the operations of each transaction are executed consecutively, without any interleaved operations
- Formally, a schedule  $S$  is *serial* if, for every transaction  $T$  participating in the schedule, all operations of  $T$  are executed consecutively
  - Otherwise the schedule is called nonserial
- In a serial schedule, only one transaction is active at a time
  - The commit (or abort) of the active transaction initiates execution of the next transaction
- An assumption that can be made is: *every serial schedule is correct*
  - All transactions should be independent (**I**solation)
  - Each transaction is assumed to be correct if executed on its own (**C**onsistency preservation)
  - Hence, the ordering of transactions in a serial schedule does not matter
    - \* Once every transaction is executed from beginning to end

## Issues with Serial Schedules

- That is all great...but:
  - Serial schedules limit concurrency
  - If a transaction is waiting for an operation to complete it is not possible to switch processing to another transaction
  - If a transaction T is very long, all other transactions must wait for T to complete its operations before they can begin
- Hence, serial schedules are considered unacceptable in practice

## Serializability

- However, if it can be determined which other schedules are *equivalent* to a serial schedule, those schedules can be allowed to occur
- How is equivalence measured?
  - Result equivalence
  - Conflict Equivalence
- If a nonserial schedule meets these criteria, it is said to be *serializable*

## Result Equivalence

- This is the simplest notion of equivalence
- Two schedules are said to be *result equivalent* if they produce the same final state of the DB
- However two schedules could coincidentally produce the same result

## Conflict Equivalence

- Two schedules are said to be *conflict equivalent* if the order of any two conflicting operations is the same in both schedules
- Reminder: two operations in a schedule are said to *conflict* if:
  1. They belong to *different transactions*
  2. They access the *same item X*
  3. And *at least one* of the operations is a write(X)
- If two conflicting operations are applied in different orders in two schedules, the effect can be different on the DB
  - Hence, the schedules are not conflict equivalent ) If read and write operations occur in the order
  - $r_1(X), w_2(X)$  in schedule  $S_1$

- $r_2(X)$ ,  $w_2(X)$  in schedule  $S_2$
  - The value read by  $r_1(X)$  may be different in the two schedules as it may have been updated by the write
- Similarly, if two write operations occur in the order:
  - $w_1(X)$ ,  $w_2(X)$  in  $S_1$
  - $w_2(X)$ ,  $w_1(X)$  in  $S_2$
  - Then the next  $r(X)$  operation in the two schedules will read potentially different values
  - If these are the last two operations in the schedules, the final value of item  $X$  in the DB will be different

## Serializability

- Being able to say that a schedule is serializable, is the same as saying it is correct
  - All serial schedules are correct
  - This is equivalent to a serial schedule
  - Hence, this is also correct
- Serializable schedules give the benefits of concurrent execution without giving up correctness
- Most DBMS systems will have a set of protocols which:
  - Must be followed by every transaction
  - Are enforced by the concurrency control subsystem
  - Ensure the serializability of all schedules in which the transactions participate
- Concurrency Control Protocols

## Concurrency Control Protocols

- Locking Protocols
  - Data items are locked to prevent multiple transactions from accessing them concurrently
- Timestamps
  - A unique identifier is generated for each transaction based upon transaction start time
- Optimistic Protocols
  - Based upon validation of the transaction after it executes its operations



## Locking

- A lock is a variable associated with a data item
  - Describes the status of the data item with respect to operations that can be applied to it
  - Data items may be at a variety of granularities, e.g. DB, table, tuple, attribute, etc.
- Locks are used to synchronise access by concurrency transactions
- There are two main types of lock
  - Binary Lock
  - Read/Write Lock

### Binary Lock

- A binary lock can have only two states (or values)
  - Locked and unlocked
- Two operations are used in binary locking
  - lock\_item
  - unlock\_item
- Each transaction locks the item being using it, and then unlocks it when finished
- A transaction which wants to access a data item requires to lock the item
  - If the item is unlocked, the transaction locks it and can use it
  - If the item is already locked, then the transaction must wait until it is unlocked
- Binary locking is rarely used, as it is too restrictive
  - At most, one transaction can access each item at a time
  - Several transactions should be allowed access concurrently if they only need read access

### Read/Write Lock

- If multiple transactions want to read an item, then they can access the item concurrently
  - Read operations are not conflicting
- However, if a transaction is to write an item, then it must have exclusive access

- The read/write lock implements this form of locking
  - It is called a multiple mode lock
- There are three locking operations used in read/write locks
  - read\_lock
  - write\_lock
  - unlock
- A *read\_locked* item is also called *share\_locked*, as other transactions are allowed to read it
- A *write\_locked* item is also called *exclusive\_locked* as a single transaction has access to it
- The following rules must be enforced
  1. A transaction *T* must issue the operation read\_lock(X) or write\_lock(X) before any read\_item(X) operation is performed in *T*
  2. A transaction *T* must issue the operation write\_lock(X) before any write\_item(X) operation is performed in *T*
  3. A transaction *T* must issue the operation unlock(X) after all read\_item(X) and write\_item(X) operations are completed in *T*
  4. A transaction *T* will not issue an unlock(X) operation unless it already holds a read (shared) lock or a write (exclusive) lock on item *X*

## Lock Conversion

- Some DBMS allow the conversion of a lock from one state to another
- A transaction *T* can issue a read\_lock(X) and then later *upgrade* the lock by issuing a write\_lock(X)
  - If *T* is the only transaction holding a read lock *X* at the time it issues the write\_lock(X) operation, the lock can be upgraded
  - Otherwise, the transaction must wait
- It is also possible for a transaction *T* to issue a write\_lock(X) and then later to *downgrade* the lock by issuing a read\_lock(X) operation

## Two-Phase Locking

- Two-Phase Locking is an additional protocol
  - Concerns the positioning of locking and unlocking in every transaction
  - Used to guarantee serializability

- A transaction is said to follow two-phase locking if all lock operations (read\_lock, write\_lock) precede the first unlock operation in that transaction
- Two-Phase Locking transactions are divided into two phases:
  - Expanding Phase
    - \* During which new locks on items can be acquired but none can be released
  - Shrinking phase
    - \* During which existing locks can be released but no new locks can be acquired
- If lock conversion is allowed
  - Upgrading of locks (from read-locked to write-locked) must be done during the expanding phase
  - Downgrading of locks (from write-locked to read-locked) must be done in the shrinking phase
- It has been proven that:
  - If every transaction in a schedule follows the two-phase locking protocol, the schedule is guaranteed to be *serializable*
- This removes the need to test each schedule produced for serializability
  - However, this comes at a cost

### Limitations

- Can limit the concurrency that can occur in a schedule
- A transaction may not be able to release an item X after it is finished using it, if the transaction must lock an additional item Y at a later point
  - Conversely, a transaction may have to lock the additional item Y before it needs it so that it can release X
  - Hence, X must remain locked until all items and the transaction needs to read or write have been locked, only then can X be released
  - Meanwhile, another transaction seeking to access X may be forced to wait

### Problems with Locking

- The use of locking in schedules can cause two additional problems
  - Deadlock
  - Starvation

## Deadlock

- Deadlock occurs when
  - Each transaction  $T_i$  in a set of two or more transactions is waiting for some item that is locked by some other transaction in the set
  - Hence, each transaction in the set is waiting for one of the other transactions in the set to release the lock on an item
  - But because the other transaction is also waiting, it will never release the lock

## Deadlock Prevention

- Used to decide what to do with a transaction involved in a possible deadlock situation:
  - Should it be blocked and made wait?
  - Should it be aborted?
  - Should the transaction preempt another transaction and cause it to abort?
- No waiting
- Wait-die
- Would-wait
- Cautious Waiting
- Typically use the concept of transaction timestamp
  - Unique identifier assigned to each transaction
- Timestamps are based on the order in which transactions are started
  - Hence, if transactions  $T_1$  starts before transactions  $T_2$ , then  $TS(T_1) < TS(T_2)$

## No Waiting

- In the *no wait* approach, if a transaction is unable to obtain a lock, it is immediately aborted
- It is then restarted after a certain time delay without checking whether a deadlock will actually occur or not
- In this case, no transaction ever waits, so no deadlock will occur
  - However, this scheme can cause transactions to abort and restart endlessly

### Wait-die

- Suppose that transaction  $T_i$  tries to lock an item  $X$  but is not able to because  $X$  is locked by some other transaction  $T_b$
- The rules followed by *Wait-die* are:
  - If  $T_a$  is older than  $T_b$  then  $T_a$  is allowed to wait
  - Otherwise, if  $T_a$  is younger than  $T_b$  then abort  $T_a$  and restart it later with the same timestamp
  - So  $T_a$  either waits or dies

### Wound-wait

- Suppose that transaction  $T_a$  tries to lock an item  $X$  but is not able to because  $X$  is locked by some other transaction  $T_b$
- The rules followed by *Wound-wait* are:
  - $T_a$  is older than  $T_b$  then abort  $T_b$  and restart it later with the same timestamp
  - Otherwise, if  $T_a$  is younger than  $T_b$  then  $T_a$  is allowed to wait
  - So  $T_a$  either wounds  $T_b$  or waits

### Wait-die and Would-wait

- Of the two transactions that may be involved in a deadlock, both these approaches abort the transaction that started later
  - Assumption is that this will waste less processing
- However, both may cause some transactions to be aborted and restarted needlessly
  - Even though those transactions may never actually cause a deadlock

### Cautious Waiting

- Proposed to try to reduce the number of needless aborts
- Suppose that transaction  $T_a$  tries to lock an item  $X$  but is not able to do so because  $X$  is locked by some other transaction  $T_b$
- The cautious waiting rules are:
  - If  $T_b$  is not blocked (not waiting for some other locked item), then  $T_a$  is blocked and allowed to wait
  - Otherwise abort  $T_a$

## Deadlock Detection

- Deadlock detection is used to check if a state of deadlock actually exists
  - Attractive is there is little interference among transactions, I.e. If different transactions rarely access the same items at the same time
  - However, if transactions are long and each transaction uses many items, it may be better to use a deadlock prevention protocol
- A simple way to detect deadlock is for the system to construct and waiting a *wait-for graph*
  - One node is created in the wait-for graph for each transaction that is currently executing
  - Whenever a transaction  $T_a$  is waiting to lock an item  $X$  that is currently locked by a transaction  $T_b$ , a directed edge ( $T_a \rightarrow T_b$ ) is created in the wait-for graph
  - When  $T_b$  releases the lock(s) on the items that  $T_a$  is waiting for, the directed edge is dropped from the wait-for graph
- There is a state of deadlock if and only if the wait-for graph has a loop
- Once detected, deadlock must be resolved by aborting and rolling back one of the transactions
- Choosing which transaction to abort is known as *victim selection*
  - Avoid selecting transactions that have run for a long time or have performed many updates
  - Instead, select transactions that have not made many changes (younger transactions)

## Starvation

- Another problem that may occur when using locking is *starvation*
  - When a transaction cannot proceed for an indefinite period of time while other transactions in the system continue normally
- Can occur if the waiting scheme for locked items gives priority to some transactions over others
- Can also occur because of victim selection
  - If the algorithm repeatedly selects the same victim transaction, causing it to abort and never finish

## Starvation Solution

- First-come-first-served queue

- Transactions are allowed to lock an item in the order in which they originally requested the lock
- Waiting list
  - Allows some transactions to have priority over others, but increases the priority of a transaction the longer it waits, until it eventually gets the highest priority and proceeds
- Victim priority
  - Assigns higher priorities to transactions that have been aborted multiple times

### Timestamp Ordering

- As already discussed, timestamp values are assigned to transactions in the order they are submitted to the system
  - So a timestamp can be thought of as the transaction start time
- There are concurrency control techniques based purely on timestamp ordering which do not use locking
  - Hence, deadlocks cannot occur
- For each data item accessed by conflicting operations in the schedule, the order in which the item is accessed must not violate the timestamp ordering
- This produces serializable schedules which are equivalent to the serial schedule
  - As all conflicting operations are conducted in the same order that they could be in the serial schedule
- In order to enforce this, the DBMS keeps two timestamp values for each data item:
  - read\_TS(X) - This is equal to the timestamp of the most recently started (youngest) transaction that has successfully read item X
  - write\_TS(X) - This is equal to the timestamp of the most recently started (youngest) transaction that has successfully written item X
- Whenever a transaction T issues a write\_item(X) operation, the following is checked:
  - If read\_TS(X) > TS(T) or if write\_TS(X) > TS(T), then abort and rollback T and reject the operation
    - \* This should be done because some younger transaction with a timestamp greater than TS(T) - and hence after T in the timestamp ordering - has already read or written the value of item X before T has a chance to write X, thus violating the timestamp ordering

- If  $\text{read\_TS}(X) \leq \text{TS}(T)$  and if  $\text{write\_TS}(X) \leq \text{TS}(T)$ , then execute the  $\text{write\_item}(X)$  operation of  $T$  and set  $\text{write\_TS}(X)$  to  $\text{TS}(T)$
- Whenever a transaction  $T$  issues a  $\text{read\_item}(X)$  operation, the following is checked
  - If  $\text{write\_TS}(X) > \text{TS}(T)$ , then abort and rollback  $T$  and reject the operation
    - \* This should be done because some younger transaction with timestamp greater than  $\text{TS}(T)$  - and hence after  $T$  in the timestamp ordering - has already written the value of item  $X$  before  $T$  has a chance to read  $X$
  - If  $\text{write\_TS}(X) \leq \text{TS}(T)$ , then execute the  $\text{read\_item}(X)$  operation of  $T$  and set  $\text{read\_TS}(X)$  equal to the larger of  $\text{TS}(T)$  and the current  $\text{read\_TS}(X)$