

Contents

Limitations of Relational Databases	1
Emergence of NoSQL	2
NoSQL is used as simple hashtag to refer to the event	2
NoSQL - No Definition	2
Distribution Models	2
Sharding	2
Replication	3
Combined Approaches	3
CAP Theorem	3
Relational vs NoSQL	4
Four Common Types of NoSQL	4
Key-Value Stores	4
Document Databases	5
Column Stores	5
Good for	5
Graph Stores	5
Graph Stores vs Relational Databases	6

Limitations of Relational Databases

- Impedance mismatch
- Application and integrations
- Scale up vs scale out

Emergence of NoSQL

- First appeared in the 90s are the name of an open-source relational database introduced by Carlo Strozzi
 - Two companies played a key role
 - 2004: Google started the project BigTable and published a paper in 2006
 - 2007: Amazon published the research paper on Amazon Dynamo
 - John Oskarsson decides to find out more about the new tendency to alternative data storage
 - 2009: Meetup in San Francisco
- **NoSQL is used as simple hashtag to refer to the event**

NoSQL - No Definition

- Non Relational
- No SQL as query language
- Schema less
- Usually - not always - Open Source Project
- They are distributed
- RDBMS use ACID transactions, NoSQL don't

Distribution Models

Sharding

- Different data on different nodes
- Each server acts as a single source for the subset of data it is responsible for
- Ideal setting: one user talks with one server
 - Data accessed together are stored together
 - Example: access based on physical location, place data to the nearest server
- Many NoSQL databases offer auto sharding
- Scales read and write on the different nodes of the same cluster
- No resilience if used alone: node failure \Rightarrow data unavailability

Replication

- The same data is replicated and copied over multiple nodes
- Master-slave: one node is the primary responsible for processing the update to data while the other are secondaries used for read operations
 - Scaling by adding slaves
 - Processing incoming data limited by master
 - Read resilience
 - Inconsistency problem
- The same data is replicated and copied over multiple nodes
- Peer-to-peer: all replicas have equal weight and can accept writing
 - Scaling by adding nodes
 - Node failure without losing write capability
 - Inconsistency problem

Combined Approaches

- Master-slave & sharding
 - Multiple master, each data has a single master
- P2P & sharding (common for column-family databases)
 - Replication of the shard

CAP Theorem

- Consistency
 - All nodes see the same data at the same time
- Availability
 - A guarantee that every request receives a response about whether it succeeded or failed
- Partition tolerance
 - The system continues to operate despite arbitrary partitioning due to network failures

Relational vs NoSQL

- Acid
 - Atomic
 - Consistent
 - Isolated
 - Durable
- Base
 - Basic Availability
 - Soft-state
 - Eventual consistency

Four Common Types of NoSQL

- Aggregate
 - Key-Value Stores
 - Document Stores
 - Column Stores
- Graph Stores
- Note
 - Lots of hybrids
 - Lots of NewSQL vendors
 - Some niche Graph stores

Key-Value Stores

- Maps keys to values
- Values treated as a blob
 - They can be complex compound objects
- Single index
- Consistency applicable for operations on a single key
- Very fast and scalable
- Inefficient to do aggregate queries, “all the carts work \$100 or more” or to represent relationships between data
- Great for shopping carts, user profiles and preferences, storing session information

Document Databases

- A document is like a hash, with one id and many values
- Store Javascript Documents
 - JSON = JavaScript Object Notation
 - An associative array
 - Key value pairs
 - Values can be documents or arrays
 - Arrays can contain documents
- Data is implicitly denormalised

Column Stores

- Store data as columns rather than rows
 - Columns organised in column family
 - Each column belongs to a single column family
 - Column acts as a unit for access
 - Particular column family will be accessed together
- Efficient to do column ordered operations
- Not so great as row based queries
- Adding columns is quite inexpensive and is done on a row-by-row basis
- Each row can have a different set of columns, or none at all, allowing tables to remain sparse without incurring a storage cost for null values

Good for

- Relational
 - Queries that return small subsets of rows
 - Queries that use a large subset of row data
- Column
 - Queries that require just a column of data
 - Queries that require a small subset of row data

Graph Stores

- Data model composed by nodes connected by edges
 - Nodes represent entities
 - Edges represent the relationships between entities

- Nodes and edges can have properties
- Querying a graph database means traversing the graph by following the relationships
- Pros
 - Representing objects of the real world that are highly inter-connected
 - Traversing the relationships in these data models is cheap

Graph Stores vs Relational Databases

- Relational databases are not ideally suited to representing relationships
- Relationships implemented through foreign keys
- Expensive joins required to navigate relationships
 - Poor performance for highly connected data models