

# Contents

<b>Transport Layer</b>	<b>3</b>
<b>Transport Services and Protocols</b>	<b>3</b>
Internet Transport-Layer Protocols . . . . .	3
<b>Multiplexing/Demultiplexing</b>	<b>4</b>
Addresses . . . . .	4
How Demultiplexing Works . . . . .	4
Connection Orientated Demux . . . . .	4
<b>User Datagram Protocol (UDP)</b>	<b>5</b>
Header . . . . .	5
<b>Principles of Reliable Data Transfer</b>	<b>6</b>
RDT . . . . .	6
rdt1.0: Reliable Transfer Over a Reliable Channel . . . . .	6
rdt2.0 . . . . .	7
rdt2.1: Discussion . . . . .	8
rdt2.2: NAK-free Protocol . . . . .	8
rdt3.0: Channels with Errors and Loss . . . . .	8
rdt3.0 . . . . .	9
<b>Pipelined Protocols</b>	<b>9</b>
Overview . . . . .	9
Go-back-N . . . . .	9
Sender . . . . .	9
Receiver . . . . .	10
Selective Repeat . . . . .	10
Sender . . . . .	10
Receiver . . . . .	10

<b>TCP (Overview)</b>	<b>11</b>
Header . . . . .	11
Seq Numbers & ACKs . . . . .	12
Round Trip Time (RTT) . . . . .	12
Average RTT . . . . .	12
Variations in RTT (DevRTT) . . . . .	13
Reliable Data Transfer . . . . .	14
Sender Events . . . . .	14
ACK Generation . . . . .	15
Fast Retransmit . . . . .	15
Flow Control . . . . .	15
<b>Connection Management</b>	<b>16</b>
TCP 3-way Handshake . . . . .	16
Opening . . . . .	16
Closing . . . . .	17
SYN Flood Attack (DoS) . . . . .	17
SYN Cookies . . . . .	17
<b>Principles of Congestion Control</b>	<b>17</b>
Causes/Costs . . . . .	18
Scenario 1 . . . . .	18
Scenario 2 . . . . .	18
Scenario 3 . . . . .	19
Approaches Towards congestion Control . . . . .	19
ATM ABR Congestion Control . . . . .	19
TCP Congestion Control Details . . . . .	20
<b>TCP</b>	<b>21</b>
Slow Start . . . . .	21
Detecting and Reacting to Loss . . . . .	21
Switching from Slow Start to CA . . . . .	21
Securing . . . . .	21

## Transport Layer

- Transport Layer Services
- Multiplexing and Demultiplexing
- Connectionless Transport: UDP
- Connection-Oriented Transport: TCP
- TCP Congestion Control

## Transport Services and Protocols

- Provide *logical communication* between app processes running on different hosts
  - Header
  - Intermediate routers don't run transport layer (Only physical, data link and network)
  - End points notice if there's any missing information
- Transport protocols run in end systems
  - Send side: breaks app messages into *segments*, passes to network layer
    - \* Packet loss would use up lots of bandwidth
    - \* Maximum Transfer Unit (MTU)
  - Rcv side: reassembles *segments* into messages, passes to app layer
- More than one transport protocol available to apps
  - Internet: TCP and UDP

## Internet Transport-Layer Protocols

- TCP - Reliable, in-order delivery
  - Congestion Control
  - Flow Control
  - File transmission
  - Connection Setup
- UDP - Unreliable, unordered delivery
  - No frills extension of “best effort” IP
  - Video calls
- Services not available
  - Delay guarantees
  - Bandwidth guarantees

## Multiplexing/Demultiplexing

- Multiplexing at sender:
  - Handle data from multiple sockets, add transport header (later used for demultiplexing)
- Demultiplexing at receiver
  - User header info to deliver received segments to correct socket

## Addresses

- MAC Address
  - Physical Layer
  - Wired into wireless card
- IP Address
  - Network Layer
- Socket Address (Port number)
  - Software Address
  - Differentiate which packet goes to which process
  - In TCP, negotiated in setup
  - In UDP, must already know

## How Demultiplexing Works

- Host receives IP datagrams
  - Each datagram has source IP address, destination IP address
  - Each datagram carries one transport layer segment
  - Each segment has source, destination port number
- Host uses *IP address & port numbers* to direct segment to appropriate socket

## Connection Orientated Demux

- TCP socket identified by 4-tuple (You need this data to find the right socket in a machine)
  - Source IP address
  - Source port numbers
  - Dest IP address

- Dest port number
- Receiver uses all four values to direct segment to appropriate socket
  - Demultiplexing
- Server host may support many simultaneous TCP sockets
  - Each socket identifier by its own 4-tuple
- Web servers have different sockets for each connecting client
  - Non-persistent HTTP will have a different socket for each request
    - \* A nonpersistent connection is the one that is closed after the server sends the requested object to the client
    - \* In other words, the connection is used exactly for one request and one response

## User Datagram Protocol (UDP)

- “No frills”, “bare bones” Internet transport protocol
- “Best effort” service, UDP segments may be
  - Lost
  - Delivered out-of-order to app
- Connectionless
  - Speed
  - Each UDP segment handled independently of others
- Users
  - Streaming multimedia apps (loss tolerant, rate sensitive)
  - DNS
  - SNMP

## Header

- No connection establishment
  - Which can add delay
- Simple - No connection state
  - At sender, receiver
- Small header size
- No congestion control
  - UDP can blast away as fast as desired
  - Can pump as much data into the network as possible




Figure 1: UDP segment format

## Principles of Reliable Data Transfer

- Important in application, transport, link layers
- Characteristics of unreliable channel will determine complexity of *reliable data transfer* protocol (rdt)
  1. **rdt\_send()**: called from above (e.g. by app), Passed data to deliver to receiver upper layer
  2. **udt\_send()**: called by rdt, to transfer packet over unreliable channel to receiver
    - Passes through unreliable channel
    - Bits might get corrupted, packets could get lost, etc.
  3. **rdt\_rcv()**: called when packet arrives on rcv-side of channel
  4. **deliver\_data()**: called by **rdt** to deliver data to upper

## RDT

- We will incrementally develop sender, receiver sides of reliable data transfer protocol
- Consider only unidirectional data transfer
  - But control info will flow in both directions!
- Use finite state machines (FSM) to specify sender, receiver

### rdt1.0: Reliable Transfer Over a Reliable Channel

- Underlying channel perfectly reliable
  - No bit errors
  - No loss of packets
- Separate FSMs for sender and receiver

- Sender sends data into underlying channel
- Receiver reads data from underlying channel

## **rdt2.0**

### **Channel with Bit Errors**

- Underlying channel may flip bits in packet
  - Checksum to detect bit errors
  - Cyclic redundancy checking (CRC) tells us bits were flipped but not which ones were flipped
- Q: How to recover from errors?
  - Acknowledgements (ACKs)
  - Receiver explicitly tells sender that pkt received OK
- Negative Acknowledgements (NAKs)
  - Receiver explicitly tells sender that pkt had errors
    - \* Sender retransmits pkt on receipt of NAK
- New mechanism in rdt2.0 (beyond rdt1.0)
  - Error detection
  - Feedback
    - \* Control msgs (ACK, NAK) from receiver to sender

### **Fatal Flaw**

- What happens if ACK/NAK corrupted
  - Sender does not know what happened to receiver!
  - Cannot just retransmit
    - \* Possible duplicate
- Handling Duplicates
  - Sender retransmits current pkt if ACK/NAK corrupted
  - Sender add sequence number to each packet
  - Receiver discards (does not deliver up) duplicate pkt
- Stop and Wait
  - Sender sends one packet, then waits for receiver response

### **rdt2.1: Discussion**

- Sender
  - seq # added to pkt
  - Two seq #'s (0, 1) will suffice - why?
  - Must check if received ACK/NAK corrupted
  - Twice as many states
    - \* State must “remember” whether “expected” pkt should have seq # of 0 or 1
- Receiver
  - Must check if received packet is duplicate
  - Note: receiver cannot tell if its last ACK/NAK received OK at sender

### **rdt2.2: NAK-free Protocol**

- Same functionality as rdt2.1
  - Using ACKs only
- Instead of NAK, receiver sends ACK for last pkt receiver OK
  - Receiver must explicitly include seq # of pkt being ACKed
- Duplicate ACK at sender results in same action as NAK
  - Retransmit current pkt

### **rdt3.0: Channels with Errors and Loss**

- New Assumption
  - Underlying channel can also lose packets (data, ACKs)
  - Checksum, seq #, ACKs, retransmissions will be of help
    - \* But not enough
- Approach
  - Retransmits if no ACK received in this time
  - If pkt (or ACK) just delayed (not lost)
    - \* Retransmission will be duplicate, but seq #'s already handle this
    - \* Receiver must specify seq # of pkt being ACKed
  - Requires countdown timer



### rdt3.0

- rdt3.0 performance stinks
  - e.g. 1 Gbps link (R), 15 ms prop. delay, 8000 bit packet (L)
- Effective throughput of 267 kbps over a 1 Gbps link

## Pipelined Protocols

- Pipelining: sender allows multiple “in-flight”, yet-to-be-acknowledged pkts
  - Range of sequence numbers must be increased
  - Buffering at sender and/or receiver
- Two generic forms of pipelined protocols
  - Go-back-N
  - Selective Repeat

### Overview

- Go-back-N (GBN)
  - Sender can have up to N unacked packets in the pipeline
  - Receiver can send cumulative acks
  - Sender has timer for oldest acked packet
    - \* When timer expires, retransmit all unacked packets
- Selective Repeat (SR)
  - Sender can have up to N unacked packets in pipeline
    - \* Receiver sends individual ack for each packet
  - Sender maintains timer for each unacked packet
    - \* When timer expires, retransmit only that unacked packet
  - Less bandwidth used in selective repeat, may be preferable over GBN

### Go-back-N

#### Sender

- k-bit seq # in pkt header
  - “Window” of up to N consecutive unacked pkts allowed
- Timer for oldest in-flight pkt
  - Timeout(n): Retransmit packet n and all higher seq # pkts that have been previously sent in window

## Receiver

- ACK-only: always send ACK for correctly-received pkt with highest in-order seq #
  - Need only remember **expected seq num**
- Out-of-order pkt
  - Discard: no receiver buffer!
    - \* Re-ACK pkt with highest in-order seq #

## Selective Repeat

- Receiver individually acknowledges all correctly received pkts
  - Buffers pkts, as needed, for eventual in-order delivery to upper layer
- Sender only resends pkts for which ACK not received
- Sender window
  - N consecutive seq #'s
  - Limits seq #s of sent, unACKed pkts

## Sender

- Data from above
  - If next available seq # in window, send pkt
- Timeout(n)
  - Resend pkt n, restart timer
- ACK(n) in [sendbase, sendbase+N]
  - Mark pkt n as received
  - If n smallest unACKed pkt (seq # == send-base), advance window base to next unACKed seq #

## Receiver

- Pkt n in [rscbase, rcvbase+N-1]
  - Send ACK(n)
  - Out-of-order: buffer
  - In-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt
- Pkt n in [rcvbase-N, rcvbase-1]

- Ack(n) - previously acknowledged pkt
- Otherwise
  - Ignore

## TCP (Overview)

- Point-to-point
  - One sender, one receiver
- Reliable, in order byte stream
  - No “message boundaries”
  - Streaming protocol
- Pipelined
  - TCP congestion and flow control set window size
- Full duplex data
  - Bi-directional data flow in same connection
  - MSS: maximum segment size
- Connection-orientated
  - Handshaking (exchange of control msgs) initialise sender, receiver state before data exchange
- Flow controlled

## Header

- Control Flags
  - URG: Urgent data (generally not used)
  - ACK: ACK # valid
  - PSH: Push data now (generally not used)
  - RST, SYN, FIN: connection estab (setup, teardown commands)
- Checksum: Internet Checksum
- Sequence/Acknowledgement number: Counting by bytes of data (not segments)
- Receive window: # bytes rcvr willing to accept

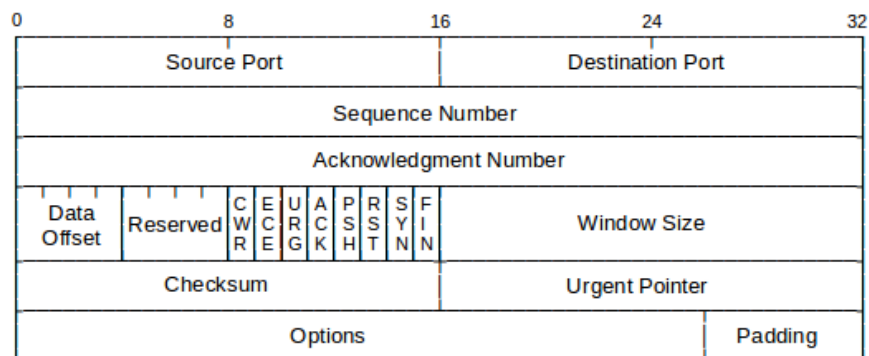


Figure 2: TCP Header

## Seq Numbers & ACKs

- Sequence Numbers
  - Byte stream “number” of first byte in segment’s data
- Acknowledgements
  - Seq # of next byte expected from other side

## Round Trip Time (RTT)

- How to set TCP timeout value to recover from lost segments?
  - Longer than RTT - but RTT varies
- How to estimate RTT?
  - SampleRTT: measured time from segment transmission until ACK receipt
    - \* Ignore retransmissions
  - SampleRTT will vary due to congestion and load at routers
    - \* Want EstimatedRTT - smooth
      - Average several recent measurements, not just current SampleRTT

## Average RTT

$$EstimatedRTT = (1 - \alpha) * EstimatedRTT + \alpha * SampleRTT$$

- Exponential weighted moving average
- Influence of past sample decreases exponentially fast
- Typical value:  $\alpha = 0.125$

Pkt #	SampleRTT	EstimatedRTT $\alpha = 0.125$
20		40
21	30	39
22	80	44
23	40	43
24	90	49
25	50	?

- What happens if  $\alpha$  is too small (say very close 0)
  - May lead to under or over estimation of RTT for a long time
- What happens if  $\alpha$  is too large (say very close 1)
  - Transient fluctuations/changes in network load affects EstimatedRTT and makes it unstable when it should not
    - \* Also leads to under or over estimation of RTT

### Variations in RTT (DevRTT)

- Timeout Interval: **EstimatedRTT** plus “safety margin”
- Estimate SampleRTT deviation from EstimatedRTT

$$DevRTT = (1 - \beta * DevRTT + \beta | SampleRTT - EstimatedRTT |)$$

(typically,  $\beta = 0.25$ )

$$TimeoutInterval = EstimatedRTT + 4 * DevRTT$$

- For TCP, initial value of TimeoutInterval = 1

Pkt #{} Pkt #	Sample RTT	EstimatedRTT $\alpha = 0.125$	DevRTT $\beta = 0.25$
20		40	10
21	30	39	10
22	80	44	18
23	40	43	14
24	90	49	20
25	50	54	25

## Reliable Data Transfer

- TCP creates rdt service on top of IP's unreliable service
  - Pipelined Segments
  - Cumulative ACKs
  - Single Retransmission Timer
- Retransmissions triggered by
  - Duplicate ACKs
- Let us initially consider simplified TCP sender
  - Ignore Duplicate ACKs
  - Ignore Flow Control & Congestion Control

## Sender Events

- Data rcvd from app
  - Create segment with seq #
- Start timer if not already running
  - Think of timer as for oldest unacked segment
    - \* Expiration interval: TimeoutInterval
- Timeout
  - Retransmit segment that caused timeout
    - \* Restart timer
- ACK rcvd
  - If ACK acknowledges previously unacked segments
    - \* Update what is known to be ACKed
    - \* Start timer if there are still unacked segments

## ACK Generation

Event At Receiver	TCP Receiver Action
Arrival of in-order segment with expected seq # already ACKed	Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK.
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediatively send single cumulative ACK, ACKing both in-order segments
Arrival of out-of-order segment higher-than-expected seq #. Gap detected	Immediately send <i>duplicate</i> ACK, indicating seq # of next expected byte
Arrival of segment that partially or completely fills gap	Immediately send ACK provided that segment starts at lower end of gap

## Fast Retransmit

- Time-out period often relatively long
  - Long delay before resending lost packet
- Detect lost segments via duplicate ACKs
  - Sender often sends many segments back-to-back
  - If segment is lost, there will likely be many duplicate ACKs
- TCP Fast Retransmit
  - If sender receives 3 duplicate ACKs (i.e., on the 4<sup>th</sup> ACK) from same data - “Triple Duplicate ACKs”
    - \* Resend unacked segment with smallest sequence number
    - \* Number of duplicates to wait for before retransmission is configurable
  - Likely that unacked segment lost, so do not wait for timeout

## Flow Control

Receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much too fast

- Receiver “advertises” free buffer space by including *rwnd* (receive window) value in TCP header of receiver-to-sender segments

- RcvBuffer size set via socket options
  - \* Typical default is 4096 bytes
- Many operating systems auto adjust the RcvBuffer

$rwnd = RcvBuffer - [LastByteRcvd - LastByteRead]$

- Sender limits amount of unacked (“in-flight”) data to receiver’s *rwnd* value
  - Guarantees receive buffer will not overflow

## Connection Management

- Before exchanging data, sender/receiver “handshake”
  - Agree to establish connection (each knowing the other willing to establish a connection)
  - Agree on connection parameters

## TCP 3-way Handshake

### Opening

1. Client State: *LISTEN* → *SYNSENT*
  - $SYNbit = 1$
  - $Seq = x$
2. Server State: *LISTEN* → *SYNRCVD*
  - $SYNBit = 1$
  - $Seq = y$
  - $ACKbit = 1$
  - $ACKnum = x + 1$
3. Client State: *SYNSEND* → *ESTAB*
  - $ACKbit = 1$
  - $ACKnum = y + 1$
  - Can piggyback user data at this point
4. Server State: *SYNRCVD* → *ESTAB*



## Closing

1. Client State:  $ESTAB \rightarrow FIN\_WAIT_1$ 
  - $FINbit = 1$
  - $Seq = x$
2. Server State:  $ESTAB \rightarrow CLOSE\_WAIT$ 
  - $ACKbit=1$
  - $ACKnum=x+1$
3. Client State  $FIN\_WAIT_1 \rightarrow FIN\_WAIT_2$
4. Server State:  $CLOSE\_WAIT \rightarrow LAST\_ACK$

## SYN Flood Attack (DoS)

- Attacker sends large number of TCP SYN segments
  - Without completing the third handshake step

## SYN Cookies

- Server does not know if the SYN segment is coming from a legitimate user
- Server creates an initial sequence number (ISN) or “cookie” from the hash of
  - Src IP addr & Port
  - Dest IP addr & Port
  - Timestamp
- Server sends SYNACK
  - Maintains no state info about corresponding to the SYN
- A legitimate client will return an ACK segment
  - Use the cookie information (ISN+1) in the ACK

## Principles of Congestion Control

- Informally:
  - “Too many sources sending too much data too fast for network to handle”
    - \* Different from flow control!
- Manifestations

- Lost packets
  - \* Buffer overflow at routers
- Long delays
  - \* Queuing in router buffers

## Causes/Costs

### Scenario 1

- Two senders, two receivers
- One router, infinite buffers
- Output link capacity:  $R$
- No retransmission
- Maximum per-connection throughput:  $R/2$
- Large queuing delays are experienced as the pkt arrival rate,  $\lambda_{in}$ , approaches capacity

### Scenario 2

- One router, *finite* buffers
  - Pkts will be dropped
- Idealization: *perfect knowledge*
  - Sender sends only when router buffers available
- Idealization: *known loss*
  - Packets can be lost, dropped at router due to full buffers
  - Sender only resends if packet known to be lost
- Realistic: *duplicates*
  - Packets can be lost, dropped at router due to full buffers
  - Sender times out prematurely, sending *two* copies, both of which are delivered
- “Costs of Congestion”
  - More work (retrans) for given “goodput”
  - Unneeded retransmissions
    - \* Link carries multiple copies of pkt
    - Decreasing goodput

### Scenario 3

- Four Senders (blue send/rcv and red send/rcv)
- Multihop Paths
- Timeout/Retransmit

Q: What happens as  $\lambda_{in}$  and  $\lambda'_{in}$  increase?

A: As red  $\lambda'_{in}$  increases, all arriving blue pkts at upper queue are dropped, blue throughput  $\rightarrow 0$

- Another “cost” of congestion
  - When packet dropped, any upstream transmission capacity used for that packet is wasted

### Approaches Towards congestion Control

- End-to-End Congestion Control
  - No explicit feedback from network
  - Congestion inferred from end-system observed loss, delay
  - Approach taken by TCP
- Network-Assisted Congestion Control
  - Routers provide feedback to end systems
  - Single bit indicating congestion (ATM - Asynchronous Transfer Mode)
  - Explicit rate for sender to send at

### ATM ABR Congestion Control

ABR: Available Bit Rate

- “Elastic Service”
- If sender’s path is underloaded
  - Sender should use available bandwidth
- If sender’s path is congested
  - Sender throttled to minimum guaranteed rate

RM (Resource Management) Cells

- Sent by sender, interspersed with data cells

- Bits in RM cell set by switches (“network assisted”)
  - NI bit: no increase in rate (mild congestion)
  - CI bit: Congestion indicated
- RM cells returned to sender by receiver, with bits intact

Two-byte ER (explicit rate) field in RM cell

- Congested switch may lower ER value in cell

EFCI (explicit forward congestion indication) bit in data cells: set to 1 in congested switch

- If data cell preceding RM cell has EFCI set, receiver sets CI bit in returned RM cell

### Additive Increase Multiplicative Decrease (AIMD)

- Sender increases transmission rate (window size)
  - Probing for usable bandwidth, until loss occurs
- Additive Increase
  - Increase congestion window (*cwnd*) by 1 MSS every RTT until loss detected
- Multiplicative Decrease
  - Cut *cwnd* in half after loss

### TCP Congestion Control Details

- Sender limits transmission

$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$

- *cwnd* is dynamic, function of perceived network congestion

TCP sending rate

- Send *cwnd* bytes
- Wait RTT for ACKs
- Then send more bytes

# TCP

## Slow Start

- When connection begins, increase rate exponentially until first loss event
  - Initially  $cwnd = 1$  MSS
  - Double  $cwnd$  every RTT
  - Done by incrementing  $cwnd$  for every ACK received
- Summary: initial rate is slow but ramps up exponentially fast

## Detecting and Reacting to Loss

- Loss indicated by timeout
  - $cwnd$  set to 1 MSS
- Window then grows exponentially (as in slow start) to threshold ( $ssthresh$ )
  - Then grows linearly
- Loss indicated by 3 duplicate ACKs: TCP Reno
  - Dup ACKs indicate network capable of delivering some segments
  - $cwnd$  is cut in half window then grows linearly
- TCP Tahoe always set  $cwnd$  to 1
  - Timeout or 3 duplicate ACKs

## Switching from Slow Start to CA

**Q:** When should the exponential increase switch to linear?

**A:** When  $cwnd$  gets to 1/2 of its value before timeout

- Implementation
  - Variable  $ssthresh$

## Securing

- TCP and UDP
  - No encryption
- SSL

- Provides encrypted TCP connection
  - Data integrity
- SSL is at app layer
  - Apps use SSL libraries which talk to TCP
- SSL socket API
  - Cleartext passwords sent into socket traverse internet encrypted