

Exercise or relax, for $\gamma = 0.9$

Recall (probability, reward)-matrices for exercise, relax

exercise	fit	unfit
fit	.99, 8	.01, 8
unfit	.2, 0	.8, 0

relax	fit	unfit
fit	.7, 10	.3, 10
unfit	0, 5	1, 5

Exercise or relax, for $\gamma = 0.9$

Recall (probability, reward)-matrices for exercise, relax

exercise	fit	unfit	relax	fit	unfit
fit	.99, 8	.01, 8	fit	.7, 10	.3, 10
unfit	.2, 0	.8, 0	unfit	0, 5	1, 5

$$q_0(s, a) := p(s, a, \text{fit})r(s, a, \text{fit}) + p(s, a, \text{unfit})r(s, a, \text{unfit})$$

$$V_n(s) := \max(q_n(s, \text{exercise}), q_n(s, \text{relax}))$$

$$q_{n+1}(s, a) := q_0(s, a) + .9(p(s, a, \text{fit})V_n(\text{fit}) + p(s, a, \text{unfit})V_n(\text{unfit}))$$

Exercise or relax, for $\gamma = 0.9$

Recall (probability, reward)-matrices for exercise, relax

exercise	fit	unfit
fit	.99, 8	.01, 8
unfit	.2, 0	.8, 0

relax	fit	unfit
fit	.7, 10	.3, 10
unfit	0, 5	1, 5

$$q_0(s, a) := p(s, a, \text{fit})r(s, a, \text{fit}) + p(s, a, \text{unfit})r(s, a, \text{unfit})$$

$$V_n(s) := \max(q_n(s, \text{exercise}), q_n(s, \text{relax}))$$

$$q_{n+1}(s, a) := q_0(s, a) + .9(p(s, a, \text{fit})V_n(\text{fit}) + p(s, a, \text{unfit})V_n(\text{unfit}))$$

	exercise	relax	π
fit	8	10	relax
unfit	0	5	relax

Exercise or relax, for $\gamma = 0.9$

Recall (probability, reward)-matrices for exercise, relax

exercise	fit	unfit	relax	fit	unfit
fit	.99, 8	.01, 8	fit	.7, 10	.3, 10
unfit	.2, 0	.8, 0	unfit	0, 5	1, 5

$$q_0(s, a) := p(s, a, \text{fit})r(s, a, \text{fit}) + p(s, a, \text{unfit})r(s, a, \text{unfit})$$

$$V_n(s) := \max(q_n(s, \text{exercise}), q_n(s, \text{relax}))$$

$$q_{n+1}(s, a) := q_0(s, a) + .9(p(s, a, \text{fit})V_n(\text{fit}) + p(s, a, \text{unfit})V_n(\text{unfit}))$$

	exercise	relax	π
fit	8, 16.955	10, 17.65	relax, relax
unfit	0, 5.4	5, 9.5	relax, relax

Exercise or relax, for $\gamma = 0.9$

Recall (probability, reward)-matrices for exercise, relax

exercise	fit	unfit	relax	fit	unfit
fit	.99, 8	.01, 8	fit	.7, 10	.3, 10
unfit	.2, 0	.8, 0	unfit	0, 5	1, 5

$$q_0(s, a) := p(s, a, \text{fit})r(s, a, \text{fit}) + p(s, a, \text{unfit})r(s, a, \text{unfit})$$

$$V_n(s) := \max(q_n(s, \text{exercise}), q_n(s, \text{relax}))$$

$$q_{n+1}(s, a) := q_0(s, a) + .9(p(s, a, \text{fit})V_n(\text{fit}) + p(s, a, \text{unfit})V_n(\text{unfit}))$$

	exercise	relax	π
fit	8, 16.955, 23.812	10, 17.65, 23.685	relax, relax, exercise
unfit	0, 5.4, 10.017	5, 9.5, 13.55	relax, relax, relax

Temporal difference (TD)

A sequence of values

$$v_1, v_2, v_3, \dots$$

averages at time k to

$$A_k := \frac{v_1 + \dots + v_k}{k}$$

Temporal difference (TD)

A sequence of values

$$v_1, v_2, v_3, \dots$$

averages at time k to

$$A_k := \frac{v_1 + \dots + v_k}{k}$$

which learning v_{k+1} updates to

$$\begin{aligned} A_{k+1} &= \frac{v_1 + \dots + v_k + v_{k+1}}{k+1} \\ &= \frac{k}{k+1} A_k + \frac{1}{k+1} v_{k+1} \end{aligned}$$

Temporal difference (TD)

A sequence of values

$$v_1, v_2, v_3, \dots$$

averages at time k to

$$A_k := \frac{v_1 + \dots + v_k}{k}$$

which learning v_{k+1} updates to

$$\begin{aligned} A_{k+1} &= \frac{v_1 + \dots + v_k + v_{k+1}}{k+1} \\ &= \frac{k}{k+1} A_k + \frac{1}{k+1} v_{k+1} \end{aligned}$$

so if $\alpha_k = \frac{1}{k}$,

$$\begin{aligned} A_{k+1} &= (1 - \alpha_{k+1}) A_k + \alpha_{k+1} \overbrace{v_{k+1}}^{\text{new}} \\ &= \underbrace{A_k}_{\text{old}} + \alpha_{k+1} \underbrace{(v_{k+1} - A_k)}_{\text{temp diff: new-old}} \end{aligned}$$

Q-Learning

Assume v_{k+1} is derived from r_{k+1}, s_{k+1} , observed sequentially

$$s_1 \xrightarrow{a_1} r_2, s_2 \xrightarrow{a_2} r_3, s_3 \xrightarrow{a_3} \cdots \underbrace{s_k \xrightarrow{a_k} r_{k+1}, s_{k+1}}_{\text{experience from which we learn}} \xrightarrow{a_{k+1}} \cdots$$

experience from which we learn

$$v_{k+1} := r_{k+1} + \gamma \max_a Q_k(s_{k+1}, a)$$

Q-Learning

Assume v_{k+1} is derived from r_{k+1}, s_{k+1} , observed sequentially

$$s_1 \xrightarrow{a_1} r_2, s_2 \xrightarrow{a_2} r_3, s_3 \xrightarrow{a_3} \cdots \underbrace{s_k \xrightarrow{a_k} r_{k+1}, s_{k+1}}_{\text{experience from which we learn}} \xrightarrow{a_{k+1}} \cdots$$

experience from which we learn

$$v_{k+1} := r_{k+1} + \gamma \max_a Q_k(s_{k+1}, a)$$

given $0 \leq \gamma < 1$, $Q_1 : (S \times A) \rightarrow \mathbb{R}$ and $v_1 \in \mathbb{R}$, with

$$Q_{k+1}(s_k, a_k) := (1 - \alpha)Q_k(s_k, a_k) + \alpha v_{k+1}$$

for $0 \leq \alpha \leq 1$,

Q-Learning

Assume v_{k+1} is derived from r_{k+1}, s_{k+1} , observed sequentially

$$s_1 \xrightarrow{a_1} r_2, s_2 \xrightarrow{a_2} r_3, s_3 \xrightarrow{a_3} \cdots \underbrace{s_k \xrightarrow{a_k} r_{k+1}, s_{k+1}}_{\text{experience from which we learn}} \xrightarrow{a_{k+1}} \cdots$$

experience from which we learn

$$v_{k+1} := r_{k+1} + \gamma \max_a Q_k(s_{k+1}, a)$$

given $0 \leq \gamma < 1$, $Q_1 : (S \times A) \rightarrow \mathbb{R}$ and $v_1 \in \mathbb{R}$, with

$$Q_{k+1}(s_k, a_k) := (1 - \alpha)Q_k(s_k, a_k) + \alpha v_{k+1}$$

for $0 \leq \alpha \leq 1$, smelling like

$$A_{k+1} = (1 - \alpha_{k+1})A_k + \alpha_{k+1}v_{k+1} \quad \text{for } \alpha_{k+1} = \frac{1}{k+1}$$

from previous slide (on TD).

Averaging?

$$\begin{aligned} v_{k+1} &= r_{k+1} + \gamma \max_a Q_k(s_{k+1}, a) \\ \underbrace{Q_{k+1}(s_k, a_k)}_{A_{k+1}} &= (1 - \alpha) \underbrace{Q_k(s_k, a_k)}_{\neq Q_k(s_{k-1}, a_{k-1}) = A_k} + \alpha v_{k+1} \end{aligned}$$

Averaging?

$$\begin{aligned} v_{k+1} &= r_{k+1} + \gamma \max_a Q_k(s_{k+1}, a) \\ \underbrace{Q_{k+1}(s_k, a_k)}_{A_{k+1}} &= (1 - \alpha) \underbrace{Q_k(s_k, a_k)}_{\neq Q_k(s_{k-1}, a_{k-1}) = A_k} + \alpha v_{k+1} \end{aligned}$$

for a deterministic MDP

i.e., $p(s, a, s') \in \{0, 1\}$ for all s, a, s'

let $\alpha = 1$ as v_{k+1} may look-ahead further than Q_k for the experience $s_k, a_k, r_{k+1}, s_{k+1}$ (determined by s_k, a_k)

Averaging?

$$\begin{aligned}v_{k+1} &= r_{k+1} + \gamma \max_a Q_k(s_{k+1}, a) \\ \underbrace{Q_{k+1}(s_k, a_k)}_{A_{k+1}} &= (1 - \alpha) \underbrace{Q_k(s_k, a_k)}_{\neq Q_k(s_{k-1}, a_{k-1}) = A_k} + \alpha v_{k+1}\end{aligned}$$

for a deterministic MDP

i.e., $p(s, a, s') \in \{0, 1\}$ for all s, a, s'

let $\alpha = 1$ as v_{k+1} may look-ahead further than Q_k for the experience $s_k, a_k, r_{k+1}, s_{k+1}$ (determined by s_k, a_k)

for $0 < p(s, a, s') < 1$, sample s' at frequency $\propto p(s, a, s')$ to average Q as a whole (not just $Q(s, a)$ at a particular (s, a)), converging to optimal Q under certain assumptions, including

$$\sum \alpha_k = \infty \quad \text{and} \quad \sum \alpha_k^2 < \infty \quad (\text{e.g. } \alpha_k = \frac{1}{k})$$

MDP, one experience at a time

Update $q : (S \times A) \rightarrow \mathbb{R}$ via p, r for

$$q'(s, a) := \sum_{s'} p(s, a, s') (r(s, a, s') + \gamma \max_{a'} q(s', a'))$$

or pointwise via experience $s_1 \xrightarrow{a_1} r_2, s_2$ for

$$q'(s, a) := \begin{cases} \alpha(r_2 + \gamma \max_{a'} q(s_2, a')) \\ \quad + (1 - \alpha)q(s, a) & \text{if } s = s_1 \text{ and } a = a_1 \\ q(s, a) & \text{otherwise.} \end{cases}$$

MDP, one experience at a time

Update $q : (S \times A) \rightarrow \mathbb{R}$ via p, r for

$$q'(s, a) := \sum_{s'} p(s, a, s') (r(s, a, s') + \gamma \max_{a'} q(s', a'))$$

or pointwise via experience $s_1 \xrightarrow{a_1} r_2, s_2$ for

$$q'(s, a) := \begin{cases} \alpha(r_2 + \gamma \max_{a'} q(s_2, a')) \\ \quad + (1 - \alpha)q(s, a) & \text{if } s = s_1 \text{ and } a = a_1 \\ q(s, a) & \text{otherwise.} \end{cases}$$

To converge to MDP's optimal Q-value, visit every state-action pair (s, a) repeatedly (for $s \xrightarrow{a} r', s'$ with diff s', r' under p, r).

MDP, one experience at a time

Update $q : (S \times A) \rightarrow \mathbb{R}$ via p, r for

$$q'(s, a) := \sum_{s'} p(s, a, s') (r(s, a, s') + \gamma \max_{a'} q(s', a'))$$

or pointwise via experience $s_1 \xrightarrow{a_1} r_2, s_2$ for

$$q'(s, a) := \begin{cases} \alpha(r_2 + \gamma \max_{a'} q(s_2, a')) \\ \quad + (1 - \alpha)q(s, a) & \text{if } s = s_1 \text{ and } a = a_1 \\ q(s, a) & \text{otherwise.} \end{cases}$$

To converge to MDP's optimal Q-value, visit every state-action pair (s, a) repeatedly (for $s \xrightarrow{a} r', s'$ with diff s', r' under p, r).

End **episode**

$$s_1 \xrightarrow{a_1} r_2, s_2 \xrightarrow{a_2} r_3, s_3 \xrightarrow{a_3} \dots \xrightarrow{a_{n-1}} r_n, s_n$$

at an absorbing state s_n with $r(s_n, a, s_n) = 0$ for every action a .

Exploration-exploitation tradeoff

$s \xrightarrow{a} r', s'$ r', s' from environment, but a ?

$$Q_{n+1}(s, a) := \alpha[r' + \gamma \max_{a'} Q_n(s', a')] + (1 - \alpha)Q_n(s, a)$$

from functional policy $\pi : S \rightarrow A$ [e.g. $\pi_Q(s) = \arg \max_a Q(s, a)$]

Exploration-exploitation tradeoff

$s \xrightarrow{a} r', s'$ r', s' from environment, but a ?

$$Q_{n+1}(s, a) := \alpha[r' + \gamma \max_{a'} Q_n(s', a')] + (1 - \alpha)Q_n(s, a)$$

from functional policy $\pi : S \rightarrow A$ [e.g. $\pi_Q(s) = \arg \max_a Q(s, a)$]

to $\pi : (S \times A) \rightarrow [0, 1]$ s.t. $\sum_{a \in A} \pi(s, a) = 1$ for each $s \in S$

e.g. for n actions, m having $\max Q(s, \cdot)$

$$\pi_Q^\epsilon(s, a) = \begin{cases} \frac{1-\epsilon}{m} + \frac{\epsilon}{n} & \text{if } Q(s, a) \text{ is max } (\dagger) \\ \frac{\epsilon}{n} & \text{otherwise } (\ddagger) \end{cases}$$

(\dagger) says exploit: use what we know

(\ddagger) says explore: try something new (for the future)

Exploration-exploitation tradeoff

$s \xrightarrow{a} r', s' \xrightarrow{a'} \dots$ r', s' from environment, but a ?

$$Q_{n+1}(s, a) := \alpha[r' + \gamma \max_{a'} Q_n(s', a')] + (1 - \alpha)Q_n(s, a)$$

from functional policy $\pi : S \rightarrow A$ [e.g. $\pi_Q(s) = \arg \max_a Q(s, a)$]

to $\pi : (S \times A) \rightarrow [0, 1]$ s.t. $\sum_{a \in A} \pi(s, a) = 1$ for each $s \in S$

e.g. for n actions, m having $\max Q(s, \cdot)$

$$\pi_Q^\epsilon(s, a) = \begin{cases} \frac{1-\epsilon}{m} + \frac{\epsilon}{n} & \text{if } Q(s, a) \text{ is max} \quad (\dagger) \\ \frac{\epsilon}{n} & \text{otherwise} \quad (\ddagger) \end{cases}$$

(\dagger) says exploit: use what we know

(\ddagger) says explore: try something new (for the future)

SARSA: replace $\arg \max$ by policy in use

$$Q_{n+1}(s, a) := \alpha[r' + \gamma Q_n(s', a')] + (1 - \alpha)Q_n(s, a)$$