

3BA26 : Concurrent Systems I: SIMD II

David Gregg

Department of Computer Science
Trinity College Dublin

April 22, 2009

"Newton's method, also called the Newton-Raphson method, is a root-finding algorithm that uses the first few terms of the Taylor series of a function $f(x)$ in the vicinity of a suspected root. Newton's method is sometimes also known as Newton's iteration, although in this work the latter term is reserved to the application of Newton's method for computing square roots."

- <http://mathworld.wolfram.com/NewtonsMethod.html>

The Newton-Rhapson Reciprocal

One Iteration of the Newton Rhapson-Method is all we need to increase precision for the reciprocal intrinsic. Note that this is still faster than using a divide!

$$rcp_nr(x) = 2 * \frac{1}{x} - (\frac{1}{x} * (x * \frac{1}{x}))$$

SSE NR Reciprocal

```
__m128 rcp_nr(const __m128 &a)
{
    const __m128 r = _mm_rcp_ps(a);
    return _mm_sub_ps(_mm_add_ps(r, r),
        _mm_mul_ps(_mm_mul_ps(r, a), r));
}
```

NR Reciprocal SQRT

$$\frac{1}{2} * rsqrtps(x) * (3 - x * rsqrtps(x) * rsqrtps(x))$$

```
__m128 rsqrt_nr(const __m128 &a)
{
    const __m128 half = _mm_set1_ps(0.5f);
    const __m128 three = _mm_set1_ps(3.0f);

    const __m128 r = _mm_rsqrt_ps(a);
    return _mm_mul_ps(_mm_mul_ps(half, r),
        _mm_sub_ps(three,
            _mm_mul_ps(_mm_mul_ps(a, r), Ra0)));
}
```

Operating on bits

SSE provides a variety of bitwise operations that can be used to manipulate individual bits within a 128 bit vector value.

For example, it is possible to do a bitwise and of all 128 bits in two vector words:

```
__m128 a = _mm_set_ps(0.0f, 1.0f, 2.0f, 3.0f);  
__m128 b = _mm_set_ps(3.0f, 2.0f, 1.0f, 0.0f);  
__m128 c;  
  
c = _mm_and_ps(a, mask);
```

Some Bitwise Operations

<code>__m128 _mm_and_ps(__m128 a, __m128 b)</code>	<code>r = a and b</code>
<code>__m128 _mm_or_ps(__m128 a, __m128 b)</code>	<code>r = a or b</code>
<code>__m128 _mm_andnot_ps(__m128 a, __m128 b)</code>	<code>r = not a and b</code>
<code>__m128 _mm_xor_ps(__m128 a, __m128 b)</code>	<code>r = a xor b</code>

Comparisons

SSE allows us to compare 4 values at a time against 4 other values.

a =

99	88	77	66
----	----	----	----

b =

88	77	66	55
----	----	----	----

$a > b$ is obviously true in this case, but what if...

b =

88	99	66	55
----	----	----	----

Some Comparison Operations

<code>__m128 _mm_cmpeq_ps (__m128 a, __m128 b)</code>	<code>=</code>
<code>__m128 _mm_cmplt_ps (__m128 a, __m128 b)</code>	<code><</code>
<code>__m128 _mm_cmple_ps (__m128 a, __m128 b)</code>	<code>≤</code>
<code>__m128 _mm_cmpgt_ps (__m128 a, __m128 b)</code>	<code>></code>
<code>__m128 _mm_cmpge_ps (__m128 a, __m128 b)</code>	<code>≥</code>
<code>__m128 _mm_cmpneq_ps (__m128 a, __m128 b)</code>	<code>!=</code>
<code>__m128 _mm_cmpnlt_ps (__m128 a, __m128 b)</code>	<code>!<</code>
<code>__m128 _mm_cmpnle_ps (__m128 a, __m128 b)</code>	<code>!≤</code>
<code>__m128 _mm_cmpngt_ps (__m128 a, __m128 b)</code>	<code>!></code>
<code>__m128 _mm_cmpnge_ps (__m128 a, __m128 b)</code>	<code>!≥</code>

Masks

Comparison instructions return a bitmask indicating which of the constituent parts of the SSE register passed and which failed. In the previously listed instructions we have four results returned, packed into an `__m128`. So that we can write code such as:

```
if( a > b ) do_a(); else do_b();
```

SSE provides the ability to convert the `__m128` mask into a 4 bit integer using the `_mm_movemask_ps` intrinsic.

Bitmask	Meaning
1111	Comparison True for all 4 floats
0000	Comparison False for all 4 floats
1100	Comparison True for first two floats
1010	Comparison True for first and third floats

Example Usage

```
__m128 a = _mm_set1_ps(0.0f);
__m128 b = _mm_set1_ps(1.0f);

__m128 r = _mm_cmpgt_ps(a, b);

if( _mm_movemask_ps(r) == 0xF)
    printf("a is greater than b\n");
else if (_mm_movemask_ps(r) == 0)
    printf("a is NOT greater than b");
else
    printf("mixed result");
```

More direct masking

Comparison instructions return a bitmask indicating which of the constituent parts of the SSE register passed and which failed.

It is possible to use this mask directly rather than using the `_mm_movemask_ps` intrinsic.

The result of a comparison is four values, one for each of the numbers compared. If the comparison is false then the result is zero. If the comparison is true, then the result is minus one.

Note that minus one in two's complement is represented by all the bits being set to one. So the result of a comparison is all the bits in the result being set to zero, or all the bits set to one.

More direct masking

So that we can write code such as:

```
__m128 a = _mm_set_ps(0.0f, 1.0f, 2.0f, 3.0f);
__m128 b = _mm_set_ps(3.0f, 2.0f, 1.0f, 0.0f);
__m128 c, mask;

mask = _mm_cmpgt_ps(a, b);
c = _mm_and_ps(a, mask);
```

The variable `c` now contains those numbers from `a` that are greater than the corresponding numbers in `(b)`.

Problem: Write a function `max(a, b)` which returns a vector containing the maximum values of `a` and `b`.

Max and Min

SSE also supports max and min operations directly:

```
__m128 a = _mm_set_ps(0.0f, 1.0f, 2.0f, 3.0f);
__m128 b = _mm_set_ps(3.0f, 2.0f, 1.0f, 0.0f);
__m128 max, min;

max = _mm_max_ps(a, b);
min = _mm_min_ps(a, b);
```

Using these operations it is possible to build quite complex bigger functions.

SISD Example

To properly harness the power of SIMD we may need to re-order our data so that it can be more efficiently loaded into registers. For example, we have four separate data streams of floats, each loaded from a separate file. We must add up the floats in each file yielding 4 separate results.

```
float total1, total2, total3, total4;  
float *data1, *data2, *data3, *data4;
```

```
for(i=0; i<SIZE; i++){  
    total1 += data1[i];  
    total2 += data2[i];  
    total3 += data3[i];  
    total4 += data4[i];  
}
```

Possible Solution

An obvious solution might be to do the following to use SIMD instructions to speed up the loop.

```
__m128 totals;
float *data1, *data2, *data3, *data4;

for(i=0; i<SIZE; i++){
    __m128 v = _mm_setr_ps(data1[i], data2[i],
                           data3[i], data4[i]);
    totals = _mm_add_ps(totals, v);
}
```


Reordered Solution

data1 =

d1_1	d1_2	d1_3	d1_4	...
------	------	------	------	-----

data2 =

d2_1	d2_2	d2_3	d2_4	...
------	------	------	------	-----

data3 =

d3_1	d3_2	d3_3	d3_4	...
------	------	------	------	-----

data4 =

d4_1	d4_2	d4_3	d4_4	...
------	------	------	------	-----

Reordered Data

data =

d1_1	d2_1	d3_1	d4_1	d1_2	d2_2	d3_2	d4_2
------	------	------	------	------	------	------	------

Reordered Solution

By reordering our data we can get it into a SIMD register more efficiently

```
__m128 totals;  
float *data;  
  
for(i=0; i<SIZE; i+=4) {  
    __m128 v = _mm_load_ps(&data[i]);  
    totals = _mm_add_ps(totals, v);  
}
```

Shuffle

SSE provides a shuffle instruction that can be used to reorder the four values within a vector word.

It actually operates on two separate vector words and takes two 32-bit values from each of them.

```
__m128 a = _mm_set_ps(0.0, 1.0, 2.0, 3.0);
__m128 b = _mm_set_ps(4.0, 5.0, 6.0, 7.0);
__m128 c;
```

```
c = _mm_shuffle_ps(a, b, _MM_SHUFFLE(1, 0, 3, 2));
```

```
/* c now has value {2.0, 3.0, 4.0, 5.0}
```

Horizontal Operations

The operations we have used so far are horizontal operations such as

1	2	3	4
---	---	---	---

+

1	2	3	4
---	---	---	---

=

2	4	6	8
---	---	---	---

Occasionally we may find it useful to be able to perform an operation across all constituent values in a single vector.

a =

2	4	6	8
---	---	---	---

horizontal_add(a) = 20:

```
c = mm_hadd ps (a, b)
```

$$b = \begin{bmatrix} b_0 & b_1 & b_2 & b_3 \end{bmatrix}$$
$$a = \begin{bmatrix} b_0 + b_1 & b_2 + b_3 & a_0 + a_1 & a_2 + a_3 \end{bmatrix}$$

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡

Scalar Operations

The horizontal operations we have used so far have all been *Packed Operations*. These operations operate on all constituents of the vector. *Scalar Operations* operate only on the least significant portion of the vector.

$$c = a + b$$

$$a = \begin{array}{|c|c|c|c|} \hline a0 & a1 & a2 & a3 \\ \hline \end{array}$$

$$b = \begin{array}{|c|c|c|c|} \hline b0 & b1 & b2 & b3 \\ \hline \end{array}$$

After the operation, c contains the following

$$c = \begin{array}{|c|c|c|c|} \hline a0 & a1 & a2 & a3 + b3 \\ \hline \end{array}$$

Scalar Instructions

By convention SSE Packed Instructions are suffixed with `_ps`. Scalar Instructions are suffixed with `_ss`. Many of the instructions we've encountered so far have a scalar version. Check the *xmmintrin.h* header file for a full list