

A Key Principle

- ▶ Haskell execution replaces sub-expressions, by ones defined to be equal (but hopefully simpler).
- ▶ This is an example of a general principle that is very desirable in functional languages — *Referential Transparency*.
- ▶ A language is *Referentially Transparent* if
 - ▶ replacing an expression by another equal expression does not change the meaning/value of the program as a whole.
 - ▶ e.g Given program `2 * sum (3:2:1:[]) + x`, then the following are all equivalent programs:
`2 * (3 + sum (2:1:[])) + x`
`2 * (3 + 2 + 1 + 0) + x`
`2 * 6 + x`
`12 + x`

Referential Transparency (Examples)

- ▶ Referentially Transparent:
 - ▶ A function whose output depends only on its inputs.
 - ▶ Expressions built from standard arithmetic operators.
 - ▶ None of the above have any “side-effects”.
- ▶ Referentially Opaque:
 - ▶ A function whose value depends on some global variable or state-component.
 - ▶ A procedure/function that modifies global state.
 - ▶ The assignment statement.
 - ▶ A function that performs I/O, it depends on the global state of “real world”, and modifies it.
 - ▶ Most of the above are examples of “side-effects”

Why Referential Transparency matters

- ▶ Reasoning about program behaviour is easier “substituting equals for equals”
- ▶ Code optimization is much simpler
- ▶ Scope for code optimization is much greater
- ▶ A programming language where every construct is referentially transparent, *w.r.t. to the “obvious” semantics*, is called “*pure*”
 - ▶ Haskell (and Clean) are pure functional languages
 - ▶ ML, Scheme, LISP are generally considered impure functional languages (they have explicit assignment and I/O side-effects), but this is *w.r.t. a simple functional semantics* for such languages.

What Referential Transparency isn't

- ▶ Referential Transparency does *not* mean:
 - ▶ the language is functional
 - ▶ the language has no side-effects
- ▶ Referential Transparency is a property relating a language and its semantics
 - ▶ most languages can be given a semantics that makes them referentially transparent.
 - ▶ The issue is one of degree: such a semantics may be very complex.
 - ▶ **Pure** functional languages are referentially transparent *w.r.t. a relatively simple and obvious semantics*.
 - ▶ An imperative language with a *full* semantics is also referentially transparent.

a relatively simple and obvious semantics???

According to Amr Sabry, a **purely functional** language is one that:

1. is a *conservative extension* of the simply typed lambda-calculus,
2. has well-defined *call-by-value*, *call-by-need*, and *call-by-name* evaluation functions,
3. and all three evaluation functions are *weakly equivalent*.

All will be a little clearer when we see lambda-calculus after Study Week!

P.S. He thinks the notion of “referential transparency” is broken — he has a point!

Further Reading

- ▶ From [haskell.org](http://www.haskell.org/haskellwiki/Referential_transparency):
http://www.haskell.org/haskellwiki/Referential_transparency
- ▶ Linked to by the above:
<http://www.cas.mcmaster.ca/~kahl/reftrans.html>
- ▶ See <http://stackoverflow.com/questions/210835/what-is-referential-transparency> for an interesting discussion.

What's really going on?

- ▶ Haskell as a re-write system makes sense, but ...
- ▶ ... how is the rewrite system implemented ?
- ▶ We know what purity is, but now we need to understand how it is achieved.
- ▶ We need to drill down further into the execution model for Haskell

Abstract Syntax Trees

- ▶ The Haskell Parser converts Haskell source-text into internal abstract syntax trees (AST).
- ▶ These trees are built from boxes of various types and edges (pointers).
- ▶ We shall describe an execution model that manipulates these trees directly.


AST Boxes

- ▶ Atomic Values and Variables: 3, True, 'c' v




—a variable box holds its **name**, not its value!

- ▶ Data Constructors: [], :



The “cons” box has 2 pointers to relevant components


- ▶ Function Application:




The “apply” box also has 2 pointers to relevant components

AST Examples (I)

- ▶ `sum []`




- ▶ `sum (x:xs)`



AST Examples (II)

- ▶ `x + sum xs`



- ▶ Note how binary application has a “spine” of 2 @-nodes.
- ▶ Remember that Haskell functions are *partially evaluated* so
`a + b = ((+) a) b`

Application Example

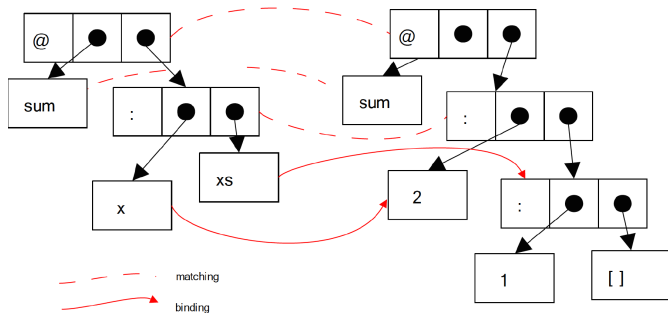
- ▶ Consider application `sum (2:1:[])`
- ▶ The LHS of `sum` is: `sum (x:xs)`
 we match this with `sum (2:1:[])`
 and bind `x ↦ 2, xs ↦ 1:[]`
 - ▶ this is done by matching the *syntax trees* recursively
 - ▶ the bindings are pointers to relevant AST fragments
- ▶ We want to replace the LHS by the RHS: `x + sum xs`,
 using the bindings above to get `2 + sum (1:[])`
 - ▶ We use RHS as a *template*,
 - ▶ we build a *copy*, replacing arguments with their bound values,
 - ▶ we replace the function application with the *copy*.
- ▶ The fact we build a *copy* of the RHS AST is crucial for referential transparency

AST Matching

- ▶ A successful match using ASTs

pattern: `sum (x:xs)`


candidate application: `sum (2:1:[])`



AST Binding


- ▶ The bindings from that successful match:

binding: `x ↦ 2, xs ↦ 1: []`




AST Copying

- ▶ The RHS from that successful match: `x + sum xs`



- ▶ The *copy* built replacing pattern variables by their bindings:

copy: `2 + sum (1:[])`




AST Shorthand

- ▶ The AST Box diagrams take up a lot of space
Let's introduce a shorthand version

- ▶ drop single boxes for basic values: `1 [] True v`
- ▶ drop triple boxes for application and cons-ing:



- ▶ So for example, `x + sum xs` now looks like:




Haskell AST Execution — another example

- ▶ The function `map` is defined as follows:

```
map f [] = []  
map f (x:xs) = (f x) : map f xs
```

- ▶ We have the following RHS ASTs:



Map AST Example

Consider application `map inc (1:2:3[])` where `inc x = x+1`


1. We match 2nd case $f \mapsto \text{inc}, x \mapsto 1, xs \mapsto 2:3:[]$
We build a *copy* of 2nd RHS, using bindings
`(inc 1) : (map inc (2:3:[]))`
2. We match 2nd case, $f \mapsto \text{inc}, x \mapsto 2, xs \mapsto 3:[]$
We build a *copy* of 2nd RHS, using bindings
`(inc 1) : ((inc 2) : (map inc (3:[])))`
3. We match 2nd case, $f \mapsto \text{inc}, x \mapsto 3, xs \mapsto []$
We build a *copy* of 2nd RHS, using bindings
`(inc 1) : ((inc 2) : ((inc 3) : (map inc [])))`
4. We match 1st case, $f \mapsto \text{inc}$
We build a *copy* of 2nd RHS, using bindings
`(inc 1) : ((inc 2) : ((inc 3) : []))`

The Importance of Copying (I)

- ▶ We clearly need to copy the function RHS, otherwise we couldn't re-use that function, because we'd have modified the definition.
- ▶ But in the application `map inc [1..3]` we not only copied the RHS, but that built us a *copy* of the original argument list.
- ▶ Couldn't a smart implementation realise that the copies simply had the leaves changed from `x` to `inc x`, and change these in place (so-called "destructive update")?

Before evaluating `map inc [1..3]`


- ▶ We have the application as an AST (simplified)



- ▶ • denotes a pointer to the expression `map inc [1,2,3]` from whatever contains that expression.
- ▶ (We show the original list as one lump)

How Haskell does `map inc [1..3]`


- ▶ We build a copy, and swing our pointer to indicate that copy



- ▶ the original list and other arguments are still present
- ▶ If there are no further pointers to the original list it becomes garbage, which is handled behind the scenes.

How we might optimise(?) `map inc [1..3]`

- ▶ We might suggest that we update the list in place and swing our application pointer to indicate that update:



- ▶ we don't alter the `map` RHS ASTs
- ▶ We (the compiler) somehow manage to see that the list structure is unchanged so we do destructive update in place.

The Importance of Copying (II)

- ▶ Destructive Update breaks Referential Transparency (w.r.t. the “natural” semantics of Haskell: programs are functions.)
- ▶ Consider the following program:

```
myfun xs = (xs, map inc xs)
```

We have paired together references to both the original `xs`, and the result of mapping `inc` across it.

- ▶ If we use copying, then the two lists returned by `myfun` are different
- ▶ If we use destructive update, then the two lists returned by `myfun` are equal.
 - ▶ but this means that `xs` and `map inc xs` are the same, which is clearly wrong.

Copying as a show-stopper (I)

- ▶ Imagine that `bigds` is a very large datastructure and `bigmod` is a function with parameters that performs large changes to it
- ▶ Copying means that the following sequence of calls is very expensive to run:

```
let bigds1 = bigmod p1 bigds
    bigds2 = bigmod p2 bigds1
    ...
    bigdsn = bigmod pn bigdsn
in ...
```

- ▶ So pure functional languages are not good for implementing large databases, processing large amounts of data, supporting design of large artifacts (i.e VLSI chips), ... ?

Copying as a show-stopper (II)

- ▶ Assume `fwrite f d` writes data `d` to file named `f` and returns a status value and `fread f` returns data read from file named `f`
- ▶ We cannot use copying to implement the following behaviour:

```
let d1 = fread "in1.dat"
    s2 = fwrite "out1.dat" (myfun d1)
    d3 = fread "out1.dat"
    s4 = fwrite "out1.dat" (another fun d3)
in ...
```

(think multipass compiler...)

- ▶ Why not? Because `fwrite` modifies the file-system of the computer on which it runs — and we *cannot* copy that !
- ▶ The side-effects in I/O are inherent, and we cannot “implement them away”.

Copying and Real-World I/O are inconsistent

- ▶ We cannot implement real-world I/O in a functionally pure language (referentially transparent w.r.t. function semantics)
- ▶ So pure functional languages are just intellectual toys ...
- ▶ Real-world functional languages (e.g. ML, Lisp, Scheme) are impure so they can
 - ▶ support real-world I/O
 - ▶ allow destructive update for large datastructures
- ▶ This slide summarises a view of (pure) functional languages still widely believed today
- ▶ This view was justifiable, until the early 1990s (Yes, that long ago !)
 - ▶ But the slide title is still correct ...