

## The Aggregate Update Problem

- ▶ With pure functional languages we are lumbered with the *aggregate update* problem:
  - ▶ manipulating large datastructures by copying is too expensive
  - ▶ performing I/O with copying is impossible
- ▶ The conclusion is that
  - ▶ We cannot allow arbitrary pure programs to do I/O
  - ▶ Any large aggregate updates will be very expensive
  - ▶ Why use the word “arbitrary” above ?

## Real-world programs aren't arbitrary

- ▶ A “real-world” programmer wouldn't write something like `(bigds, modify p bigds)`
- ▶ The idea of wanting to have both the original and modified versions of a large data-structure simultaneously available would seem absurd.
- ▶ Nor would any reasonable programmer want to access new and old versions of files.
- ▶ In practise, in real programs, once we modify something large or external, we never expect to see the old version.
  - ▶ — unless we make explicit provision for some form of “undo”
  - ▶ — or are required to maintain an “audit trail” !

## Single-Threadedness

- ▶ The use of a data value `ds` is “single-threaded”<sup>1</sup> if
    - ▶ There is only ever one live reference to it.
    - ▶ Once a function has been applied to it (`f ds`), the program no longer refers to `ds`.
  - ▶ Non-single threaded examples:
    - ▶ `f ds = (g ds, h ds)`  
We have two live references to `ds` once `f` is evaluated
    - ▶ `let ds1 = f ds`  
    `ds2 = g ds in ...`
- After `ds1`, we still proceed to refer to `ds` in the next line
- ▶ Single-threaded example:  
`let ds1 = f ds`  
    `ds2 = g ds1 in ...`

The application of `g` is to the *result* of the application of `f`

---


<sup>1</sup>Nothing to do with “threads” in concurrent programming

## Optimising Single Threadedness

- ▶ If a function is applied to a single-threaded argument ...
  - ▶ destructive update does not destroy referential transparency
- ▶ Copying is only necessary to ensure purity if arguments are “multi-threaded”
- ▶ So, the destructive update optimisation we saw for `map inc xs` is valid if no further reference to `xs` exists in the program.
- ▶ We can implement I/O (destructively) in a pure language if
  - ▶ we can ensure all references to files/outside world are single-threaded.

## Revisiting optimised(?) `map inc [1..3]`

- ▶ We might suggest that we update the list in place and swing our application pointer to indicate that update:



- ▶ If we can show there are no further live pointers to either the original list `[1,2,3]`, or the topmost `@` node, then the above destructive update is safe and does not affect referential transparency.

## Enforcing Single-Threadedness

There are two ways to enforce single-threadedness.

### 1. “Invent and Verify”

- ▶ Programmers write program in their own way.
- ▶ Compiler checks for single-threadedness where required.
- ▶ Implemented in language “Clean” using the type system (uniqueness types)
- ▶ Undecidable, so valid programs may be rejected

### 2. “Correctness by Construction”

- ▶ Required single-threadedness enforced by compiler
- ▶ i.e. Special restricted sub-language for I/O
- ▶ Implemented in Haskell using abstract `IO` type and “monads”.
- ▶ Can’t write all possible single-threaded programs

## I/O in Haskell

- ▶ We are now going to explore how Haskell supports I/O with all its side-effects, whilst also maintaining referential transparency w.r.t. the “natural” functional semantics.
- ▶ First we shall look at a few file operations to note their (destructive) side-effects
- ▶ Then we shall introduce the Haskell `IO` type constructor.

## File I/O Operations : File Open/Close

### ▶ `openFile`

**Input:** pathname, opening mode (read/write)

**Effect:** modifies filesystem by creating new file

**Return Value:** handle to new file

### ▶ `hClose`

**Input:** file handle

**Effect:** closes file indicated by handle, modifying filesystem

**Return Value:** none

- ▶ Real-world items affected: filesystem, file status

- ▶ opening modes: `ReadMode`, `WriteMode`, ...

## File I/O Operations : File Put/Get

- ▶ `hPutChar`
  - Input: file handle, character
  - Effect: modifies file by appending the character
  - Return Value: none
- ▶ `hGetChar`
  - Input: file handle
  - Effect: reads character from file current-position, which is then incremented
  - Return Value: character read
- ▶ Real-world items affected: contents of and positioning in open files

## I/O in Haskell

- ▶ Functions that do I/O (as a side-effect) use a special abstract datatype: `IO a`
- ▶ Type `IO a` denotes a “value”:
  - ▶ whose evaluation produces an I/O side-effect
  - ▶ which returns a value of type `a` when evaluated.
  - ▶ such “values” are called “I/O-actions”.
- ▶ I/O-actions that don’t return a value have type `IO ()`
  - ▶ Type `()` is the singleton (a.k.a. “unit”) type
  - ▶ It has only one value, also written `()`
- ▶ I/O-actions are usually invoked using special syntax (“do-notation”).

## Other File I/O Types

- ▶ `data IOMode = ReadMode | WriteMode | ...`  
file opening mode
- ▶ `type FilePath = String`  
File Pathname — just a string.
- ▶ `data Handle = ....`  
File Handles — *pointers* to open files
- ▶ There are no types to represent file themselves, or the file-system.
- ▶ I/O in Haskell works by hiding references to external data that is destructively updated.

## Haskell File I/O Functions

- ▶ `openFile :: FilePath -> IOMode -> IO Handle`  
`openFile fp mode` is an I/O-action that opens a file and returns a new handle.
- ▶ `hClose :: Handle -> IO ()`  
`hClose f` is an I/O-action that closes file `f`, returning nothing.
- ▶ `hPutChar :: Handle -> Char -> IO ()`  
`hPutChar f c` is an I/O-action that writes `c` to file `f`, returning nothing.
- ▶ `hGetChar :: Handle -> IO Char`  
`hGetChar f` is an I/O-action that reads from file `f`, returning the character read.

## Haskell Console I/O Functions

- ▶ `putChar :: Char -> IO ()`  
`putChar c` is an I/O-action that writes `c` to `stdout`, returning nothing.
- ▶ `getChar :: IO Char`  
`getChar` is an I/O-action that reads from `stdin`, returning the character read.

## Composing I/O actions

- ▶ For this stuff to be useful we need some way of *composing* two actions together in a safe (that is, single-threaded) way.
- ▶ As ever in Haskell the solution is to come up with a function.
- ▶ Perhaps something with a type like this?  
`seqIO :: (IO a) -> (IO b) -> (IO b)`
- ▶ And we could use it like this:  
`putAB = seqIO ( putChar 'a' ) ( putChar 'b' )`

## Composing I/O actions

- ▶ Actually, this will read better if we make `seqIO` an infix function:  
`(>>) :: IO a -> IO b -> IO b`  
`infixl 1 >>`  
  
`putAB = putChar 'a' >> putChar 'b'`
- ▶ We need something more elaborate to handle IO actions that produce results
- ▶ If we pass on the result of the first action as an input to the second then we can make use of it in subsequent actions.

## Composing I/O actions

- ▶ `bindIO :: (IO a) -> (a -> IO b) -> IO b`
- ▶ Obviously, we want to make this infix as well
- ▶ `(>>=) :: (IO a) -> (a -> IO b) -> IO b`
- ▶ Easy to use this for simple compositions:  
`getput = getChar >>= putChar`
- ▶ We don't always want to make use of the result right away – one neat solution is to use a lambda abstraction:  
`swap2char  
 = getChar >>=  
 (\ c1 -> getChar >>=  
 (\ c2 -> putChar c2 >> putChar c1))`

## Building I/O actions

- ▶ Apart from the four file I/O actions just presented, we have ways to build our own.
- ▶ `return :: a -> IO a`  
`return x` is an I/O action that has no side-effect, and simply returns `x`.
- ▶ `get2str :: IO String`  
`get2str`  
    `= getChar >>=`  
        `(\ c -> getChar >>=`  
            `(\ d -> return [c,d]))`

## Syntactic sugar for I/O actions

- ▶ We have some syntactic sugar for the `>>` and `>>=` functions.
- ▶ Consider  
    `getChar >>= (\ c -> getChar >>= (\ d -> return [c,d]))`
- ▶ We can write a multiline version, dropping brackets, as  
    `getChar >>= \ c ->`  
    `getChar >>= \ d ->`  
    `return [c,d]`  
  
    Look at Haskell precedence and binding rules carefully.
- ▶ We can read this as:  
    “`getChar`, call it `c`, `getChar`, call it `d`, and return `[c,d]`”.

## Do-notation

- ▶ We have  
    `getChar >>= \ c ->`  
    `getChar >>= \ d ->`  
    `return [c,d]`
- ▶ Haskell has syntactic sugar for this: “do-notation”:  
    `do c <- getChar`  
    `d <- getChar`  
    `return [c,d]`
- ▶ Restriction: the only way to compose actions is to sequence them.

## Invoking Actions in do-notation

- ▶ If action `act a b c` returns nothing, we simply call it  
    `act a b c`
- ▶ If action `act a b c` returns a value we can either:
  - ▶ simply invoke it as is (discarding the return value)  
        `act a b c`
  - ▶ or invoke it and bind the return value to a variable  
        `x <- act a b c`
- ▶ The last action in a `do`-expression cannot bind its return value  
    `actlast a b c`  
  
    Its return value becomes that of the entire `do`-expression.

## File I/O Examples (I)

- ▶ Read character from one file, and write to another

```
fCopyChar :: FilePath -> FilePath -> IO ()
fCopyChar fromf tof
  = do ff <- openFile fromf ReadMode
      c <- hGetChar ff
      hClose ff
      tf <- openFile tof WriteMode
      hPutChar tf c
      hClose tf
```

- ▶ Notes:

- ▶ no explicit reference to filesystem
- ▶ no explicit reference to file/open file data, just a reference to the file handle pointer

## File I/O Examples (Ia)

- ▶ For comparison, here is the same function, but without using the “do” notation

```
fCopyChar fromf tof
  = openFile fromf ReadMode >>= \ ff ->
    hGetChar ff >>= \ c ->
    hClose ff >>
    openFile tof WriteMode >>= \ tf ->
    hPutChar tf c >>
    hClose tf
```

- ▶ Why “do” notation was invented!
- ▶ But note that despite its imperative appearance, it is just **function application** using `>>` and `>>=`.