

Stack and Heap

Stack is first-in-last-out (FILO) storage used to support function procedure calls

Heap is memory that is allocated and freed in blocks, accessed by pointers.

Stack Usage by Functions (I)

Consider the obvious way to do factorial in Haskell:

```
fac 0 = 1 ; fac n = n * fac (n-1)
```

We get the following evaluation of `fac 2`:

```
fac 2
= 2 * fac 1
= 2 * (1 * fac 0)
= 2 * (1 * 1)
```

The key thing to note here is that when a recursive call (`fac 1` say) terminates, there is still some more work to be done by its caller (`fac 2` in this case, which still has to multiply by `2`)

Stack Usage by `fac` (I)

Compilers use the stack to keep track of local values for function input parameters and return values.

So the call `fac 2` will create a stack “frame” as follows on top of the stack:

<code>n</code>	<code>2</code>
ret	\square_2
prior stack	

Stack Usage by `fac` (II)

It will then make a recursive call, noting that it needs to multiply the result by `n`.

<code>n</code>	<code>1</code>
ret	\square_1
<code>n</code>	<code>2</code>
ret	<code>2 * </code> \square_1
prior stack	

Stack Usage by fac (III)

The call **fac 1** will then itself make a recursive call, noting that it needs to multiply the result by **n**.

n	0
ret	\square_0
n	1
ret	$1 * \square_0$
n	2
ret	$2 * \square_1$
prior stack	

Stack Usage by fac (IV)

The call **fac 0** produces result **1** ...

n	0
ret	1
n	1
ret	$1 * \square_0$
n	2
ret	$2 * \square_1$
prior stack	

Stack Usage by fac (V)

We pop the stack and return the result

n	1
ret	$1 * 1$
n	2
ret	$2 * \square_1$
prior stack	

Stack Usage by fac (VI)

We pop the stack and return the result

n	2
ret	$2 * (1*1)$
prior stack	

Stack Usage by `fac` (VII)

We pop the stack and return the result $2*(1*1)$ to the original caller

prior stack

Stack Usage by Functions (II)

Consider the following less obvious way to do factorial in Haskell:

```
fac' n = factr 1 n
factr p 0 = p
factr p n = factr (n*p) (n-1)
```

We get the following evaluation of `fac' 2`:

```
fac' 2
= factr 1 2
= factr (2*1) 1
= factr (1*(2*1)) 0
= 1*(2*1)
```

The key thing to note here is that when a recursive call (`factr (2*1) 1` say) terminates, there is *no more work* to be done by its caller (`factr 1 2`).

Stack Usage by `factr` (I)

So the call `factr 1 2` will create a stack “frame” as follows on top of the stack:

p	1
n	2
ret	\square_2
stack for <code>fac' 2</code>	

Stack Usage by `factr` (II)

It will then make a recursive call, noting that it just returns the result.

p	$2*1$
n	1
ret	\square_1
p	1
n	2
ret	\square_1
stack for <code>fac' 2</code>	

Stack Usage by `factr` (III)

`factr (2*1) 1` will then make a recursive call, noting that it just returns the result.

p	1*(2*1)
n	0
ret	\square_0
p	2*1
n	1
ret	\square_0
p	1
n	2
ret	\square_1
stack for <code>fac' 2</code>	

Stack Usage by `factr` (IV)

`factr (1*(2*1)) 0` just returns the result, `p`, which is `1*(2*1)`.

p	1*(2*1)
n	0
ret	1*(2*1)
p	2*1
n	1
ret	\square_0
p	1
n	2
ret	\square_1
stack for <code>fac' 2</code>	

Stack Usage by `factr` (V)

We pop the stack and return the result

p	2*1
n	1
ret	1*(2*1)
p	1
n	2
ret	\square_1
stack for <code>fac' 2</code>	

Stack Usage by `factr` (VI)

We pop the stack and return the result

p	1
n	2
ret	1*(2*1)
stack for <code>fac' 2</code>	

Stack Usage by `factr` (VII)

We pop the stack and return the result to the caller `fac' 2`

stack for <code>fac' 2</code>

Tail Recursion

Function `factr` is *tail recursive*:

- ▶ If making a recursive call, it is the last thing it does.
- ▶ It simply passes on the result of the call to its own caller.
- ▶ We don't need a new stack frame for each recursive call!
- ▶ No need to track own local information once call is made.

This leads to so-called "tail-call optimisation"

Tail Call Optimisation `factr` (I)

So the call `factr 2` from `fac' 2` will create a stack "frame" as normal on top of the stack:

p	1
n	2
ret	\square_2
stack for <code>fac' 2</code>	

Tail Call Optimisation `factr` (II)

The call `factr 1 2` will compute `n*p` and `(n-1)` and will then simply update the stack items *in place*:

p	(2*1)
n	1
ret	\square_2
stack for <code>fac' 2</code>	

Tail Call Optimisation `factr` (III)

The call `factr (2*1) 1` will compute `n*p` and `(n-1)` and will then simply update the stack items *in place*:

p	1*(2*1)
n	0
ret	\square_2
stack for <code>fac' 2</code>	

Tail Call Optimisation `factr` (IV)

The call `factr (1*(2*1)) 0` will simply return `p`:

p	1*(2*1)
n	0
ret	1*(2*1)
stack for <code>fac' 2</code>	

Tail Call Optimisation `factr` (V)

We pop the stack and return the result to the caller `fac' 2`.

stack for <code>fac' 2</code>

Tail Recursion is a While Loop!

In effect we can implement tail recursion just like a while-loop from an imperative language.

```
fac' n = factr 1 n                p := 1 ;

factr p n                          while n > 0
  | n > 0                          { p := n*p ; n := n-1 } ;
  = factr (n*p) (n-1)

  | otherwise = p                  return p
```

Any decent compiler does this optimisation.

Lazy Evaluation

Haskell uses *Lazy (non-Strict) Evaluation*

- ▶ Expressions are only evaluated *when* their value is needed
- ▶ In particular, argument expressions are not evaluated before a function is applied
- ▶ We find this approach allows us to write sensible programs not possible if strict-evaluation is used.
- ▶ However, it comes at a price ...

Example: `isOdd`

- ▶ We define a function checking for 'oddness' as follows:

```
isOdd n = n `mod` 2 == 1
```

- ▶ Consider the call `isOdd (1+2)`

- ▶ A strict (non-lazy) evaluation would be as follows:

```
isOdd (1+2)
= isOdd 3           evaluate arg first,
= 3 `mod` 2 == 1    then apply function
= 1 == 1
= True
```

- ▶ A non-strict (lazy) evaluation would be as follows:

```
isOdd (1+2)
= (1+2) `mod` 2 == 1    apply function first,
= 3 `mod` 2 == 1        mod needs args evaluated
= 1 == 1
= True
```

`1+2` is only evaluated *when* `mod` needs its value to proceed.

`len` and `down`

- ▶ We have a length function `len`:

```
len xs = if null xs then 0 else 1 + len (tail xs)
```

- ▶ We have a function `down` that generates a list, counting down from its numeric argument:

```
down n = if n <= 0 then [] else n : (down (n-1))
```

For example, `down 3 = [3,2,1]`

- ▶ We shall consider pattern matching versions shortly

Strict evaluation of `len (down 1)`

```
len (down 1)
= len (if 1 <= 0 then [] else 1 : (down (1-1)))
= len (1 : (down (1-1)))
= len (1 : (down 0))
= len (1 : (if 0 <= 0 then [] else 0 : (down (0-1))))
= len (1 : [])
= if null (1 : []) then 0 else 1 + len (tail (1 : []))
= 1 + len (tail (1 : []))
= 1 + len []
= 1 + len (if null [] then 0 else 1 + len (tail []))
= 1 + 0
= 1
```

We have 11 steps

Lazy evaluation of len (down 1) (part 1)

```
len (down 1)
= if null xs1 then 0 else 1 + len (tail xs1)
  where xs1 = down 1
= if null xs1 then 0 else 1 + len (tail xs1)
  where xs1 = if 1 <= 0 then [] else 1 : (down (1-1))
= if null xs1 then 0 else 1 + len (tail xs1)
  where xs1 = 1 : (down (1-1))
= 1 + len (tail xs1) where xs1 = 1 : (down (1-1))
= 1 + ( if null xs2 then 0 else 1 + len (tail xs2)
      where xs2 = tail xs1 )
  where xs1 = 1 : (down (1-1))
```

(continued overleaf)

Lazy evaluation of len (down 1) (part 2)

```
1 + ( if null xs2 then 0 else 1 + len (tail xs2)
     where xs2 = tail xs1 )
  where xs1 = 1 : (down (1-1))
= 1 + ( if null xs2 then 0 else 1 + len (tail xs2)
     where xs2 = tail (1 : (down (1-1))) )
= 1 + ( if null xs2 then 0 else 1 + len (tail xs2)
     where xs2 = down (1-1))
= 1 + ( if null xs2 then 0 else 1 + len (tail xs2)
     where xs2 = (if (1-1) <= 0
                     then [] else (1-1) : (down ((1-1)-1))) )
= 1 + ( if null xs2 then 0 else 1 + len (tail xs2)
     where xs2 = (if 0 <= 0
                     then [] else (1-1) : (down ((1-1)-1))) )
= 1 + ( if null xs2 then 0 else 1 + len (tail xs2)
     where xs2 = [] )
= 1 + 0
= 1    12 steps, each more expensive!
```

Why the xs₁ = ... ?

- ▶ Consider the first step:

```
len (down 1)
= if null xs1 then 0 else 1 + len (tail xs1)
  where xs1 = down 1
```
- ▶ We don't evaluate `down 1` — we bind it to formal parameter `xs1`
- ▶ Parameter `xs` occurs twice, but we don't copy:

```
...down 1 ...down 1 ...
```

Instead we share the reference, indicated by the `where` clause:

```
...xs1 ...xs1 ...where xs1 = down 1
```
- ▶ Function `len` is recursive, so we get different instances of `xs` which we label as `xs1`, `xs2`, ...
- ▶ The grouping of an (unevaluated) expression (`down 1`) with a binding (`xs1 = down 1`) is called either a “closure”, or a “thunk”.
- ▶ Building thunks is a *necessary* overhead for implementing lazy evaluation.

Lazy Evaluation: the costs

- ▶ Lazy evaluation has an overhead: building thunks
- ▶ Memory consumption per reduction step is typically slightly higher
- ▶ In our examples so far:

```
isOdd (1+2)
len (down 1)
```

we needed to evaluate almost everything
- ▶ So far we have observed no advantage to lazy evaluation ...

Advantages of Laziness (I)

- Imagine we have a function definition as follows:

```
myfun carg struct1 struct2
= if f carg
  then g struct1
  else h struct2
```

where `f`, `g` and `h` are internal functions

- Consider the following call:

```
myfun val s1Expr s2Expr
```

where both `s1Expr` and `s2Expr` are very expensive to evaluate.

- With strict evaluation we would have to compute both before applying `myfun`
- With lazy evaluation we evaluate `f val`, and then only evaluate one of either `s1Expr` or `s2Expr`, and then, only if `g` or `h` requires its value.

Laziness and Pattern Matching

- Consider a pattern matching version of `len`

```
len [] = 0
len (x:xs) = 1 + len xs
```

- How is call `len aListExpression` evaluated ?
- In order to pattern match we need to know if `aListExpression` is empty, or a cons-node.
- We evaluate `aListExpression`, but only to the point where we know this difference
If it is not null, we do not evaluate the head element, or the tail list.

- e.g. if `aListExpression = map f (1:2:3:[])`, where

```
map f [] = []
map f (x:xs) = f x : map f xs
```

then we only evaluate `map f` as far as

```
f 1 : map f (2:3:[])
```

Advantages of Laziness (II)

- Prelude function `take n xs` returns the first `n` elements of `xs`

```
take 0 _ = []
take _ [] = []
take n (x:xs) = x : (take (n-1) xs)
```

- Function `from n` generates an *infinite* ascending list starting with `n`.

```
from n = n : (from (n+1))
```

- Evaluating `from n` will fail to terminate for any `n`.
- Evaluation of `take 2 (from 0)` depends on the evaluation method.

Strict Evaluation of `take 2 (from 0)`

```
take 2 (from 0)
= take 2 (0 : from 1)
= take 2 (0 : 1 : from 2)
= take 2 (0 : 1 : 2 : from 3)
= take 2 (0 : 1 : 2 : 3 : from 4)
= take 2 (0 : 1 : 2 : 3 : 4 : from 5)
= ...
```

(You get the idea ...)

Lazy Evaluation of `take 2 (from 0)`

Here we don't bother to show the closures explicitly (using `where xs1 = ...`).

```
take 2 (from 0)
= take 2 (0 : from 1)
= 0 : (take 1 (from 1))
= 0 : (take 1 (1 : from 2))
= 0 : (1 : take 0 (from 2))
= 0 : (1 : [])
```

We are done ! We only built the bit of `from 0` that we actually needed.

Evaluation Strategy and Termination

We can summarise the relationship between evaluation strategy and termination as:

- ▶ There are programs that simply do not terminate, no matter how they are evaluated
e.g. `from 0`
- ▶ There are programs that terminate if evaluated lazily, but fail to terminate if evaluated strictly
e.g. `take 2 (from 0)`
- ▶ There are programs that terminate regardless of chosen evaluation strategy
e.g. `len (down 1)`
- ▶ However, there are *no* programs that terminate if evaluated strictly, but fail to terminate if evaluated lazily.