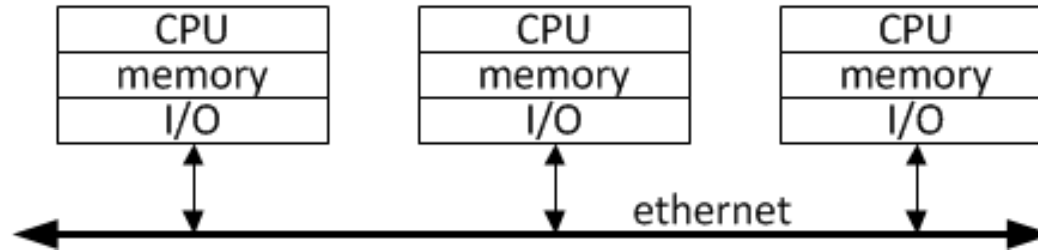


# Multiprocessors

## Loosely coupled [**Multi-computer**]



- each CPU has its own memory, I/O facilities and OS
- CPUs **DO NOT** share physical memory
- [IITAC Cluster](#) [in Lloyd building]

346 x IBM e326 compute node each with 2 x 2.4GHz 64bit AMD Opteron 250 CPUs, 4GB RAM, ...  
80GB SATA scratch disk, 2 x Broadcom BCM5704 Gbit Ethernet and PCI-X 10Gb InfiniBand connect  
2 x Voltaire 288 port InfiniBand switches + 512 port Gbit switch

Debian [**Linux**]



Theoretical peak performance 3.4TFlops [**Linpac 2.7Tflops**]

ranked 345<sup>th</sup> most powerful supercomputer [**in 2006**]

- a distributed system, ...

# Multiprocessors

## Tightly coupled [Multiprocessor]




cache coherency can be handled by CPU

- CPUs physically share memory and I/O
- inter-processor communicate via shared memory
- symmetric multiprocessing possible [i.e. single shared copy of operating system]
- critical sections protected by locks/semaphores
- processes can easily migrate between CPUs

# Multiprocessors

## Multiprocessor Cache Coherency

- what's the problem? must guarantee that a CPU always reads the most up to date value of a *memory location*



- 1) CPU A reads location X, copy of X stored in cache A
- 2) CPU B reads location X, copy of X stored in cache B
- 3) CPU A writes to location X, X in cache and main memory updated [**write-through**]
- 4) CPU B reads X, BUT gets out of date [**stale**] copy from its cache B

# Multiprocessors

## Multiprocessor Cache Coherency...

- solutions based on *bus watching* or *snoopy* caches [yet again]
- *snoopy* caches watch all bus transactions and update their internal state according to a pre-determined cache coherency protocol
- cache coherency protocols given names eg. Dragon, **Firefly**, **MESI**, Berkeley, Illinois, **Write-Once** and **Write-Through** [will discuss protocols in bold]


# Multiprocessors

## Write-Through Protocol

- key ideas
  - uses write-through caches
  - when a cache observes a bus write to a memory location it has a copy of, it simply invalidates the cache line [a write-invalidate protocol]
  - the next time the CPU accesses the same memory location/cache line, the data will be fetched from memory [hence getting the most up to date value]

# Multiprocessors

## Write-Through State Transition Diagram



- each cache line is in either the **VALID** or **INVALID** state [if an address is not in the cache, it is captured by the **INVALID** state]
- $PW \equiv$  processor write and write through to memory
- bus traffic  $\approx \Sigma$  writes [**writes  $\approx 20\%$  of memory accesses**]
- straightforward to implement and effective for small scale parallelism [**< 6 CPUs**]

# Multiprocessors

## Write-Once Protocol


- uses write-back caches [to reduce bus traffic]

BUT still uses a write-invalidate protocol

- each cache line can be in one of 4 states:
  - **INVALID**
  - **VALID**: cache line in one or more caches, cache copies  $\equiv$  memory
  - **RESERVED**: cache line in one cache ONLY, cache copy  $\equiv$  memory
  - **DIRTY**: cache line in one cache ONLY and it is the ONLY up to date copy [memory copy out of date / stale]

# Multiprocessors

## Write-Once State Transition Diagram





# Multiprocessors

## Write-Once Protocol...

- key ideas
  - write-through to VALID cache lines
  - write back (local write) to RESERVED/DIRTY cache lines [as we know cache line in one cache ONLY]
- when a memory location is read initially, it enters the cache in the VALID state
- a cache line in the VALID state changes to the RESERVED state when written for the first time [hence the name write-once]
- a write which causes a transition to the RESERVED state is written through to memory so that all other caches can observe the transaction and invalidate their copies if present [cache line now in one cache ONLY]

# Multiprocessors

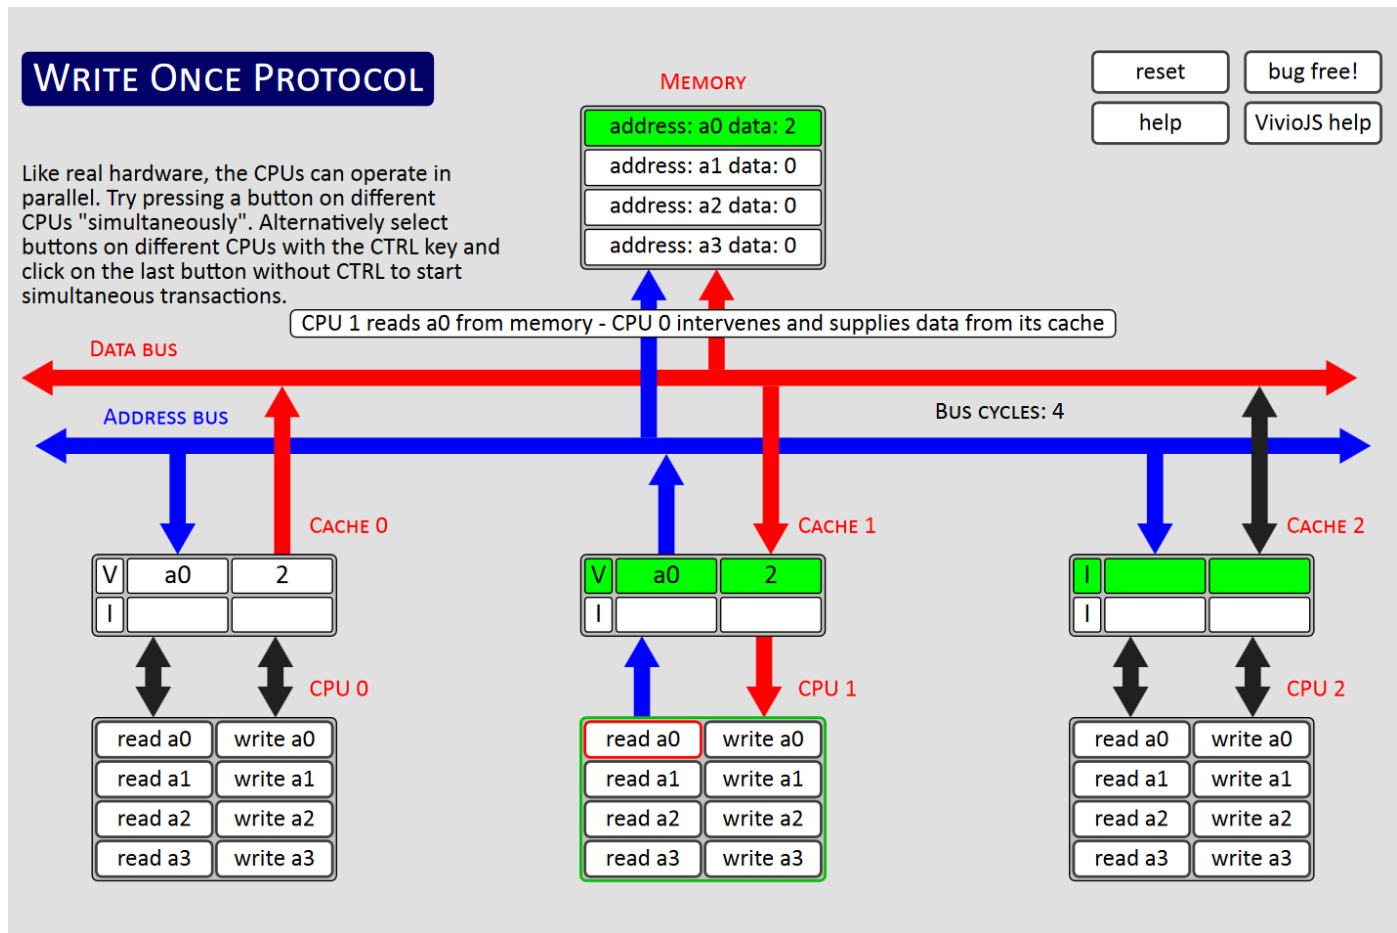
## Write-Once Protocol...

- subsequent writes to a RESERVED cache line will use write-deferred cycles and the cache line marked as DIRTY as it is now the only up to date copy in the system [memory copy out of date]
- caches must monitor bus for any reads or writes to its RESERVED and DIRTY cache lines
  - if a cache observes a read from one of its RESERVED cache lines, it simply changes its state to VALID
  - if a cache observes a read from one of its DIRTY cache lines, the cache must intervene [intervention cycle] and supply the data onto the bus to the requesting cache, the data is also written to memory and the state of both cache lines are changed to VALID
  - NB: behaviour on an observed write [DIRTY => INVALID]

# Multiprocessors

## Write-Once Protocol...

- try the Vivio cache coherency animation



# Multiprocessors

## Firefly Protocol

- used in the experimental Firefly DEC workstation

DEC  $\equiv$  Digital Equipment Corporation

- each cache line can be in one of 4 states:
  - $\sim$ Shared &  $\sim$ Dirty [Exclusive & Clean]
  - $\sim$ Shared & Dirty [Exclusive & Dirty]
  - Shared &  $\sim$ Dirty [Shared & Clean]
  - Shared & Dirty [Shared & Dirty]
- NB: there is NO INVALID state

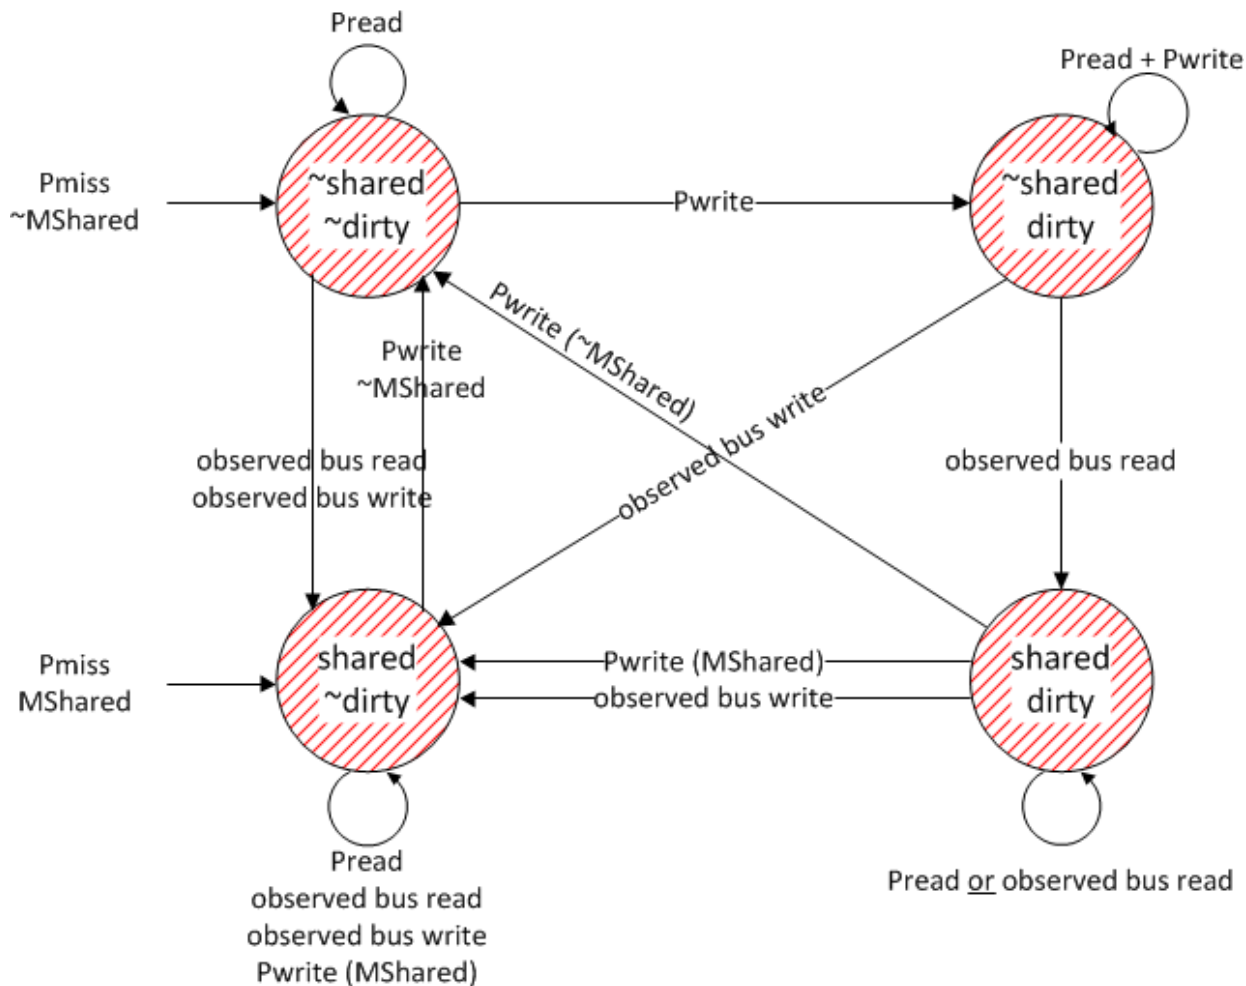
# Multiprocessors

## Firefly Protocol...

- key ideas
  - a cache knows if its cache lines are shared with other caches [**may not actually be shared but that is OK**]
  - when a cache line is read into the cache the other caches will assert a common SHARED bus line if they contain the same cache line
  - writes to exclusive cache line are write-deferred
  - writes to shared cache lines are write-through and the other caches which contain the same cache lines are updated together with memory [**write-update protocol**]
  - when a cache line ceases to be shared, it needs an extra write-through cycle to find out that the cache line is no longer shared [**SHARED signal will not be asserted**]
  - sharing may be illusory e.g. (i) processor may no longer be using a *shared* cache line  
(ii) process migration between CPUs [**sharing with an old copy of itself**]

# Multiprocessors

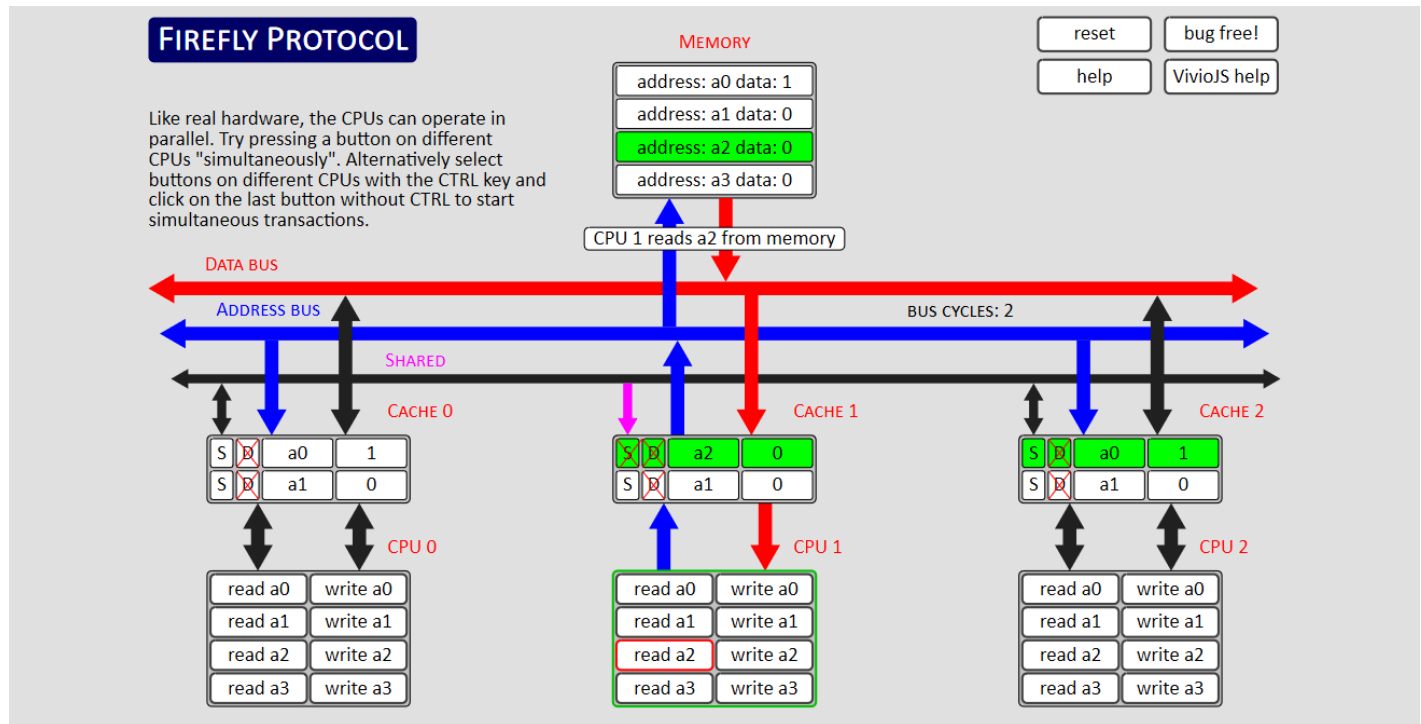
## Firefly State Transition Diagram



# Multiprocessors

## Firefly Protocol...

- try the Vivio cache coherency animation



# Multiprocessors

## Firefly Protocol...

- as there is NO INVALID state
  - at reset the cache is placed in a *miss* mode and a bootstrap program fills cache with a sequence of addresses making it consistent with memory [**VALID state**]
  - during normal operation a location can be displaced if the cache line is needed to satisfy a miss, BUT the protocol never needs to invalidate cache line
- how is the Shared & Dirty state entered?
  - asymmetrical behaviour with regard to updating memory [**avoids changing memory cycle from a read cycle to a write cycle mid cycle**]
- try the vivio animation



# Multiprocessors

## MESI Cache Coherency Protocol

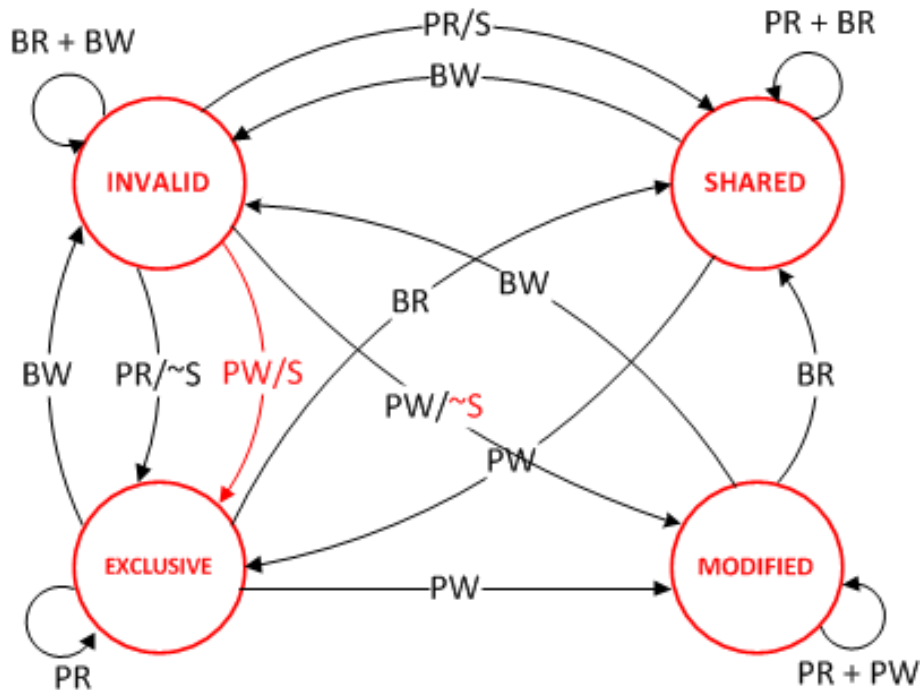
- used by Intel CPUs
- very similar to write-once, BUT uses a *shared* bus signal [**like Firefly**] so cache line can enter cache and be placed directly in the shared or exclusive state
- a write-invalidate protocol

### MESI states compared with Write-Once states

- Modified  $\equiv$  dirty
  - Exclusive  $\equiv$  reserved
  - Shared  $\equiv$  valid
  - Invalid  $\equiv$  invalid
- 
- what is the difference between the MESI and Write-Once protocols?
  - cache line may enter cache and be placed directly in the Exclusive [**reserved**] state
  - "write-once" write-through cycles no longer necessary if cache line is Exclusive
  - try the Vivio animation

# Multiprocessors

## MESI State Transition Diagram



PR = processor read  
PW = processor write

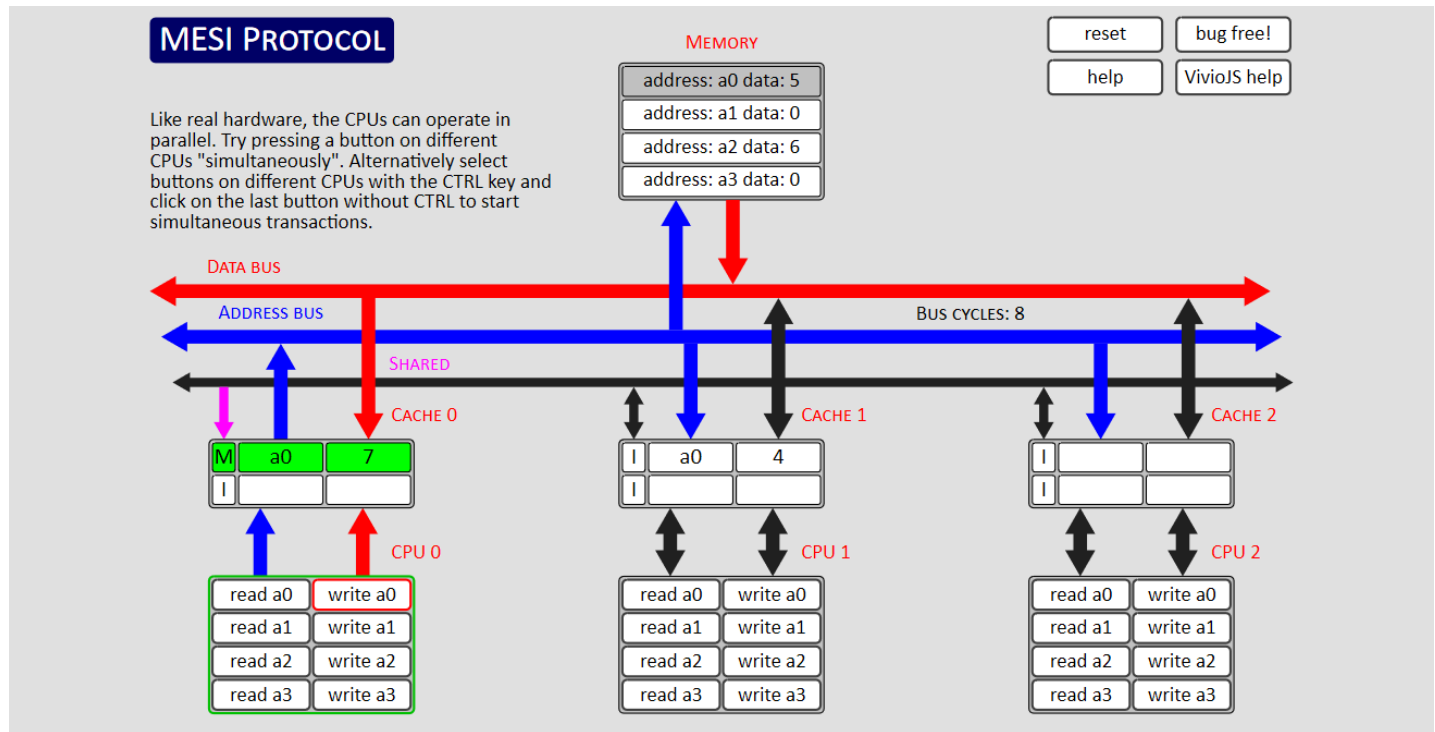
BR = observed bus read  
BW = observed bus read

S/~S = shared/NOT shared

# Multiprocessors


## MESI Protocol...

- try the Vivio cache coherency animation



# Multiprocessors

## What's next?



# Multiprocessors

## What's next?...

- network of multiprocessors
- shared bus removed as it is a bottleneck
- physical memory interleaved between CPUs
- given a physical address, it is straightforward to determine to which CPU the memory is attached
- high speed point-to-point links between CPUs [**Intel QuickPath and AMD HyperTransport**]
- if a CPU accesses non local memory, request sent via high speed point-to-point links to correct CPU which accesses its local memory and returns the data
- System can be view as a large shared memory multiprocessor
- point-to-point protocol supports cache-coherency [**Intel MESIF and AMD MOESI**]

# Multiprocessors

## Summary

- you are now able to:
  - explain the cache coherency problem
  - describe the operation of the write-through, write-once, Firefly and MESI cache coherency protocols
  - predict the cache behaviour and bus traffic for a sequence of CPU "memory" accesses in a multiprocessor with a cache coherent bus
  - evaluate the cost updating shared data