

# Overview

Recall module is roughly split into four parts:

1. Random events: counting, events, axioms of probability, Bayes, independence
2. Random variables: discrete RVs, mean and variance, correlation, conditional expectation

## Mid-term

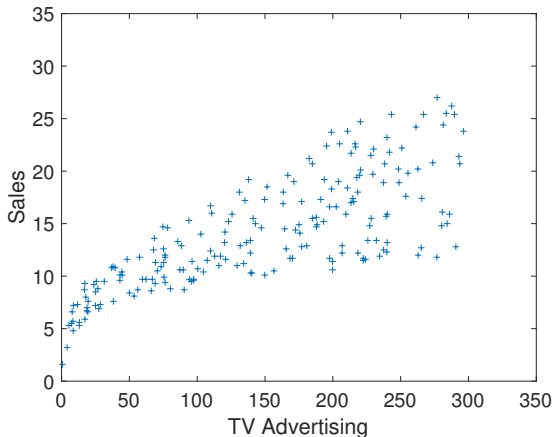
3. Inequalities and laws of large numbers: Markov, Chebyshev bounds, sample mean, weak law of large numbers, central limit theorem, confidence intervals, bootstrapping
4. Statistical models: continuous random variables, logistic regression, least squares

## Example: Advertising Data

- Data taken from An Introduction to Statistical Learning with Applications in R (<http://www-bcf.usc.edu/~gareth/ISL/data.html>)
- Data consists of the advertising budgets for three media (TV, radio and newspapers) and the overall sales in 200 different markets.

TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	9.3
⋮	⋮	⋮	⋮

## Example: Advertising Data



- Suppose we want to predict sales in a new area ?
- Predict sales when the TV advertising budget is increased to 350 ?
- ... Draw a line that fits through the data points


## Some Notation

Training data:


TV ( $x$ )	Sales ( $y$ )
230.1	22.1
44.5	10.4
17.2	9.3
$\vdots$	$\vdots$

- $m$ =number of training examples
- $x$ ="input" variable/features
- $y$ ="output" variable/"target" variable
- $(x^{(i)}, y^{(i)})$  the  $i$ th training example
- $x^{(1)} = 230.1, y^{(1)} = 22.1,$   
 $x^{(2)} = 44.5, y^{(2)} = 10.4$


# Model Representation




- Prediction:  
 $\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x$
- $\theta_0, \theta_1$  are (unknown) parameters
- sometimes abbreviate  $h_{\theta}(x)$  to  $h(x)$



$$\theta_0 = 15, \theta_1 = 0$$



$$\theta_0 = 0, \theta_1 = 0.1$$

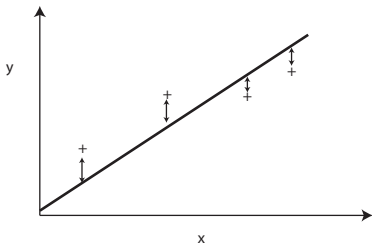
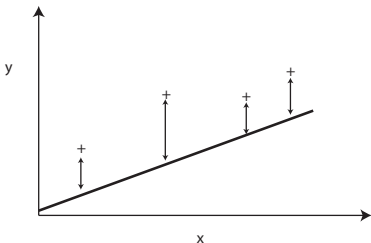


$$\theta_0 = 15, \theta_1 = 0.1$$

## Cost Function: How to choose model parameters $\theta$ ?

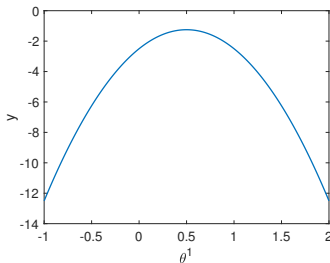
- Prediction:  $\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x$
- Idea: Choose  $\theta_0$  and  $\theta_1$  so that  $h_{\theta}(x^{(i)})$  is close to  $y^{(i)}$  for each of our training examples  $(x^{(i)}, y^{(i)})$ ,  $i = 1, \dots, m$ .
- Least squares case: select the values for  $\theta_0$  and  $\theta_1$  that minimise cost function:

$$J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$




## Simple Example

- Suppose our training data consists of just two observations:  $(1, 3)$ ,  $(2, 1)$ , and to keep things simple we know that  $\theta_0 = 0$ .
- The cost function is
$$\frac{1}{2} \sum_{j=1}^2 (y^{(j)} + \theta_1 x^{(j)})^2 = \frac{1}{2} (1 - 3\theta_1)^2 + (2 - 1\theta_1)^2$$
- What value of  $\theta_1$  minimises  $(1 - 3\theta_1)^2 + (2 - 1\theta_1)^2$  ?



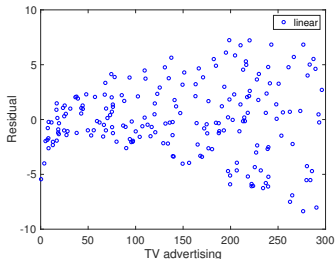
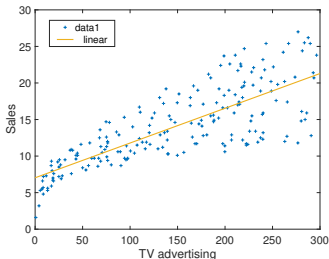
## Example: Advertising Data





## Example: Advertising Data

- Least square linear fit
- Residuals are the difference between the value predicted by the fit and the measured value.
  - Do the residuals look “random” or do they have some “structure” ?  
Is our model satisfactory ?
  - We can use the residuals to estimate a confidence interval for the prediction made by our linear fit.
- We could use cross-validation/bootstrapping to estimate out confidence in the fit itself.



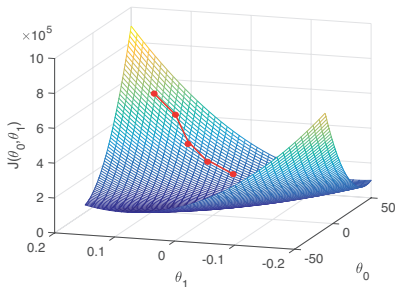
# Summary

- Hypothesis:  $h_{\theta}(x) = \theta_0 + \theta_1 x$
- Parameters:  $\theta_0, \theta_1$
- Cost Function:  $J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$
- Goal: Select  $\theta_0$  and  $\theta_1$  that minimise  $J(\theta_0, \theta_1)$

## Gradient Descent

Need to select  $\theta_0$  and  $\theta_1$  that minimise  $J(\theta_0, \theta_1)$ . Brute force search over pairs of values of  $\theta_0$  and  $\theta_1$  is inefficient, can we be smarter ?



- Start with some  $\theta_0$  and  $\theta_1$
- Repeat:
  - Update  $\theta_0$  and  $\theta_1$  to new value which makes  $J(\theta_0, \theta_1)$  smaller



- When curve is “bowl shaped” or convex then this must eventually find the minimum.

# Gradient Descent

- Start with some  $\theta_0$  and  $\theta_1$
- Repeat:
  - Update  $\theta_0$  and  $\theta_1$  to new value which makes  $J(\theta_0, \theta_1)$  smaller
- When curve has several minima then we can't be sure which we will converge to.
- Might converge to a local minimum, not the global minimum



# Gradient Descent

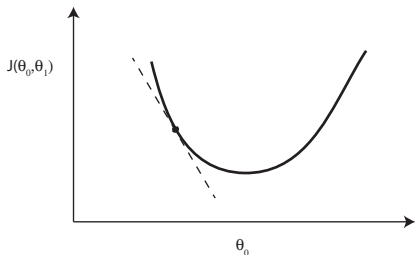
Repeat: Update  $\theta_0$  and  $\theta_1$  to new value which makes  $J(\theta_0, \theta_1)$  smaller

- One option: carry out local search of  $\theta_0$  and  $\theta_1$  to find one that decreases  $J$ .
- Another option: gradient descent:

$$temp0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

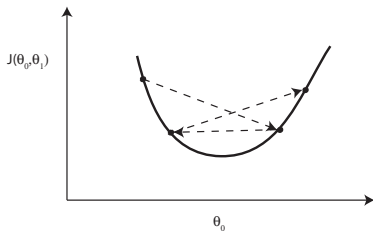
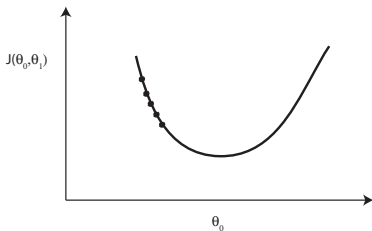
$$temp1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := temp0, \theta_1 := temp1$$



- $\frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) \approx \frac{J(\theta_0 + \delta, \theta_1) - J(\theta_0, \theta_1)}{\delta}$  for  $\delta$  sufficiently small.
- $J(\theta_0 + \delta, \theta_1) \approx J(\theta_0, \theta_1) + \delta \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$
- When  $\delta = -\alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$  then  
 $J(\theta_0 + \delta, \theta_1) \approx$   
 $J(\theta_0, \theta_1) - \alpha \left( \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) \right)^2$

# Gradient Descent



- Selecting step size  $\alpha$  too small will mean it takes a long time to converge to minimum
- But selecting  $\alpha$  too large can lead to us overshooting the minimum
- We need to adjust  $\alpha$  so that algorithm converges in a reasonable time.

## Gradient Descent

For  $J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$  with  $h_{\theta}(x) = \theta_0 + \theta_1 x$ :

- $\frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{2}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$
- $\frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{2}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x^{(i)}$

So gradient descent algorithm is:

- repeat:  
     $temp0 := \theta_0 - \frac{2\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$   
     $temp1 := \theta_1 - \frac{2\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x^{(i)}$   
     $\theta_0 := temp0, \theta_1 := temp1$

# Linear Regression with Multiple Variables

Advertising example:

TV $x_1$	Radio $x_2$	Newspaper $x_3$	Sales $y$
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	9.3
$\vdots$	$\vdots$	$\vdots$	$\vdots$

- $n$ =number of features (3 in this example)
- $(x^{(i)}, y^{(i)})$  the  $i$ th training example e.g.

$$x^{(1)} = [230.1, 37.8, 69.2]^T = \begin{bmatrix} 230.1 \\ 37.8 \\ 69.2 \end{bmatrix}$$

- $x_j^{(i)}$  is feature  $j$  in the  $i$ th training example, e.g.  $x_2^{(1)} = 37.8$



# Linear Regression with Multiple Variables

Hypothesis:  $h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$

e.g.  $h_{\theta}(x) = 15 + 0.1 \underbrace{x_1}_{\text{TV}} - 5 \underbrace{x_2}_{\text{Radio}} + 10 \underbrace{x_3}_{\text{Newspaper}}$

*Sales*                      *TV*                      *Radio*                      *Newspaper*

- For convenience, define  $x_0 = 1$   
i.e.  $x_0^{(1)} = 1, x_0^{(2)} = 1$  etc

- Feature vector  $x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$

- Parameter vector  $\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix}$

- $h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n = \theta^T x$

# Linear Regression with Multiple Variables

- Hypothesis:  $h_{\theta}(x) = \theta^T x$  (with  $\theta, x$  now  $n + 1$ -dimensional vectors)
- Cost Function:  $J(\theta_0, \theta_1, \dots, \theta_n) = J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$
- Goal: Select  $\theta$  that minimises  $J(\theta)$

As before, can find  $\theta$  using:

- Start with some  $\theta$
- Repeat:  
    Update vector  $\theta$  to new value which makes  $J(\theta)$  smaller

e.g using gradient descent:

- Start with some  $\theta$
- Repeat:  
    for  $j=0$  to  $n$   $\{ tempj := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \}$   
    for  $j=0$  to  $n$   $\{ \theta_j := tempj \}$

## Gradient Descent with Multiple Variables

For  $J(\theta) = \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$  with  $h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \dots + \theta_n x_n$ :

- $\frac{\partial}{\partial \theta_0} J(\theta) = \frac{2}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$
- $\frac{\partial}{\partial \theta_1} J(\theta) = \frac{2}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)}$
- $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{2}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$

So gradient descent algorithm is:

- Start with some  $\theta$
- Repeat:
  - for  $j=0$  to  $n$   $\{tempj := \theta_j - \frac{2\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}\}$
  - for  $j=0$  to  $n$   $\{\theta_j := tempj\}$

## Example: Advertising Data

- How is the impact of the advertising spend on TV and radio related, if at all ?
- Perhaps a quadratic fit would be better ? If so, what does that imply for how we allocate our advertising budget ?

